Pure Exploration in Multi-Armed Bandits

Zixin Zhong (University of Alberta)

Vincent Y. F. Tan (National University of Singapore)

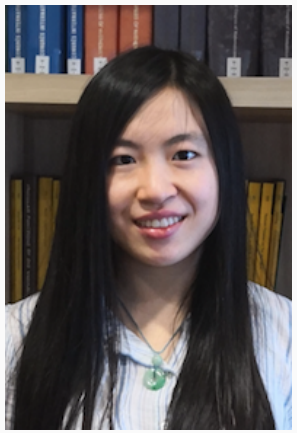

National University of Singapore
Tutorial 2 in IJCAI 2022
25 July 2022

Outline

1. **What is multi-armed bandits (MAB)?**
   - Classification of MAB problems
   - Example — Cascading bandits

Outline

1 What is multi-armed bandits (MAB)?

Motivation: data-driven optimization

- Subdomain of reinforcement learning, online learning problem.
- Application:
  - Internet advertisement placement
  - Restaurant recommendation
  - Clinical trials
  - .....

Style of tutorial:

- Will present a few well-known models/algorithms
- Will present some "newer" models/algorithms
- Since it's a tutorial, we will go through some proofs

Multi-armed bandit problem (MAB)

Multi-armed bandit problem (MAB)

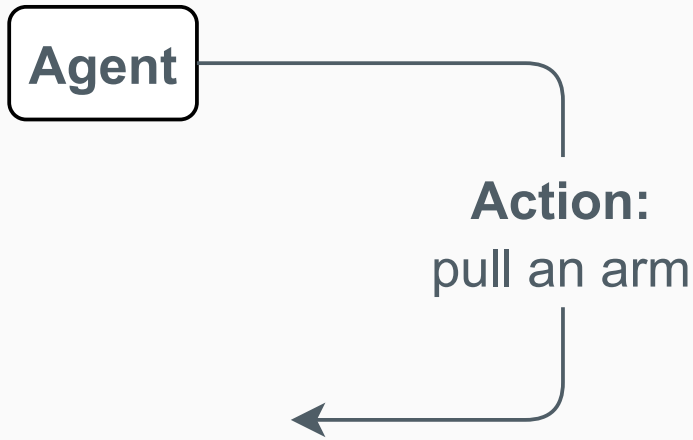

Multi-armed bandit problem (MAB)

Multi-armed bandit problem (MAB)

Objectives

1. Maximize the cumulative reward over a fixed horizon.
2. Find the best arm (largest expected reward).

Multi-armed Bandit problem (MAB)

Multi-armed Bandit problem (MAB)

Challenge

- Exploitation: to pull "confident" arms to maximize reward.
- Exploration: to pull "unconfident" arms to find better ones.

Outline

1 **What is multi-armed bandits (MAB)?**

- Classification of MAB problems

Classification

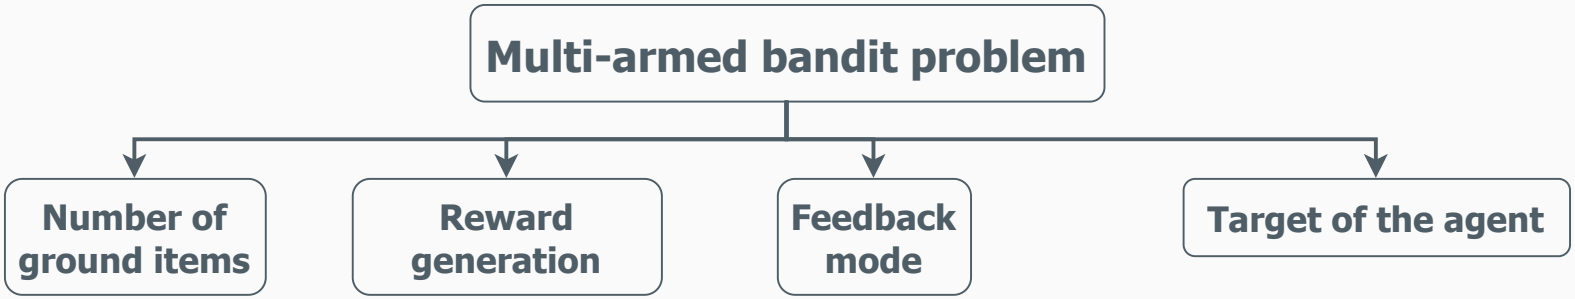

Classification

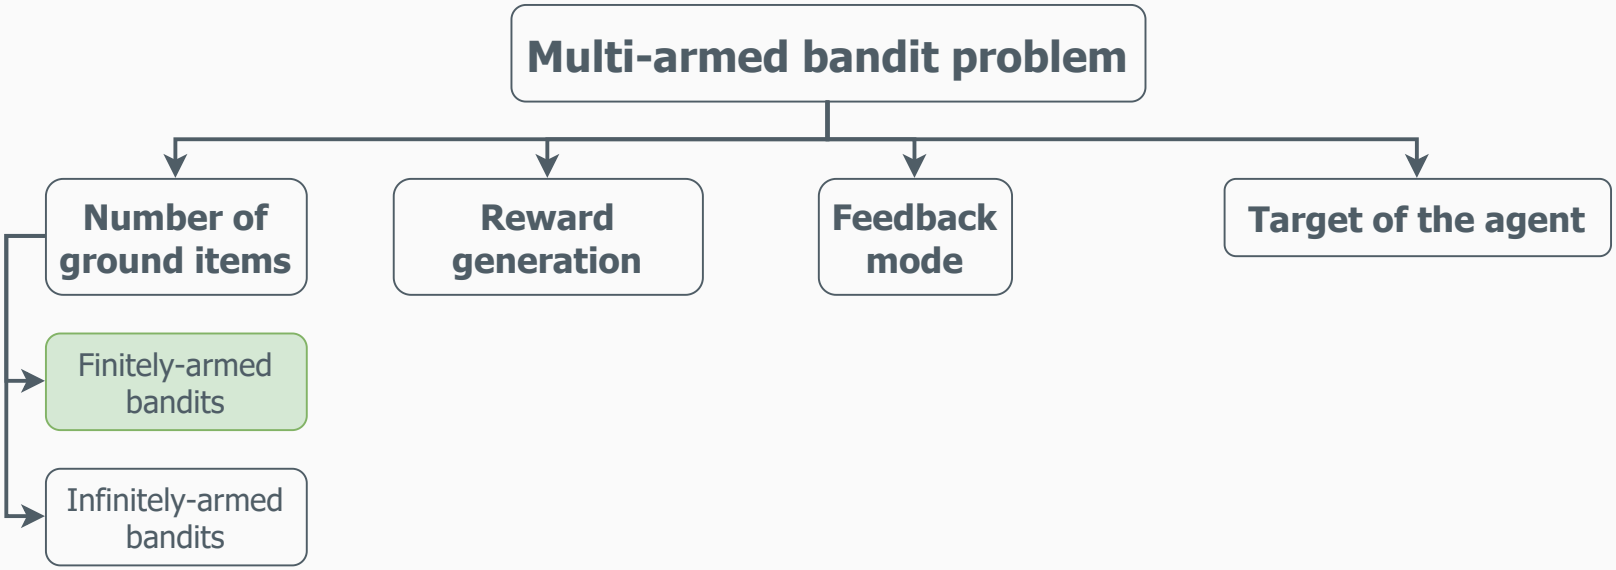

Classification

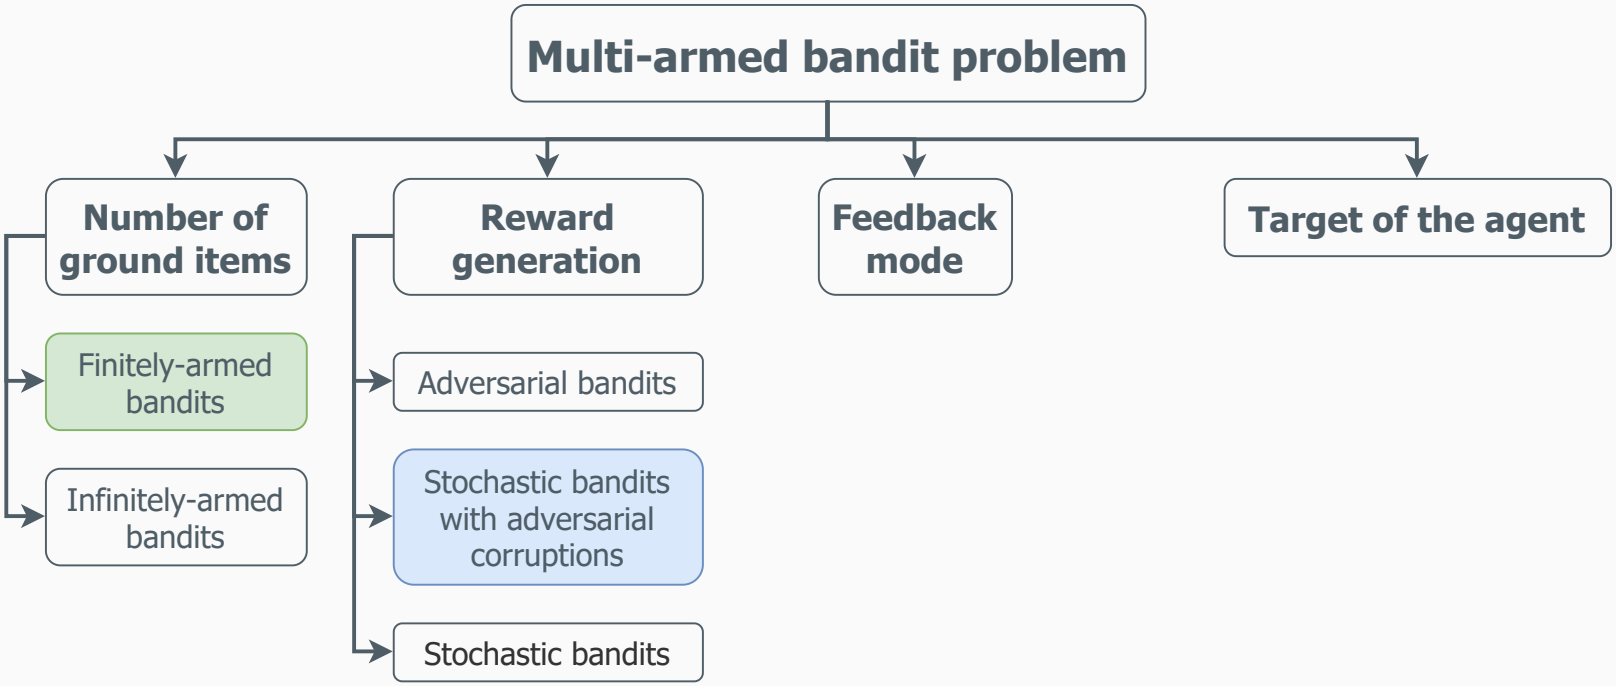

Classification

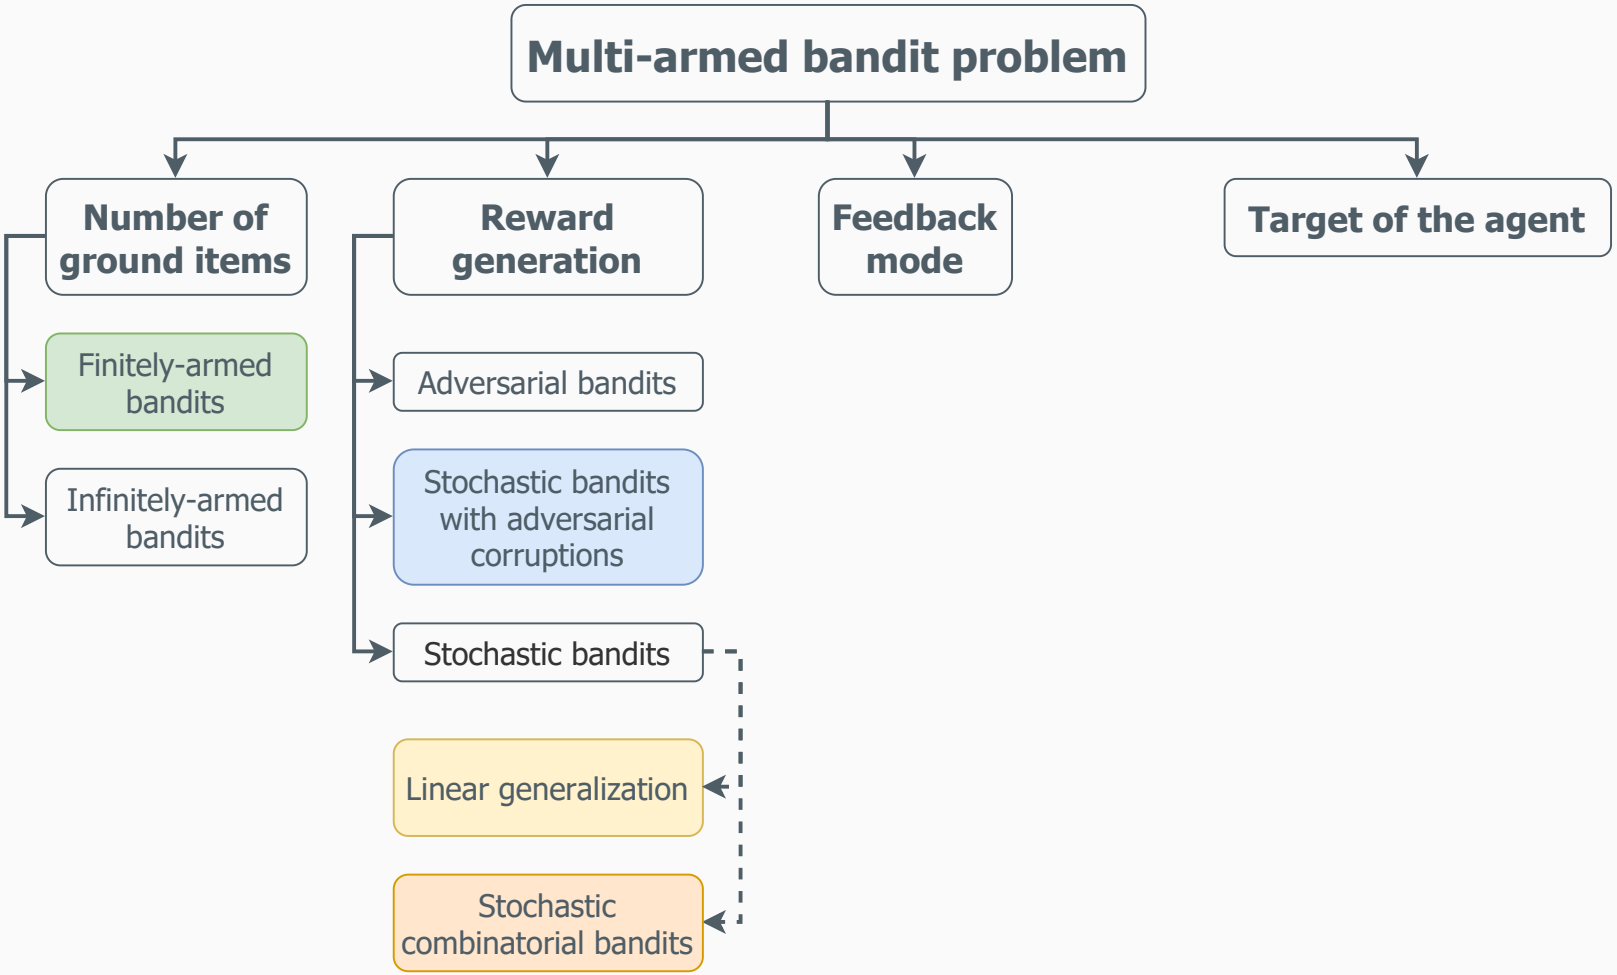

Classification

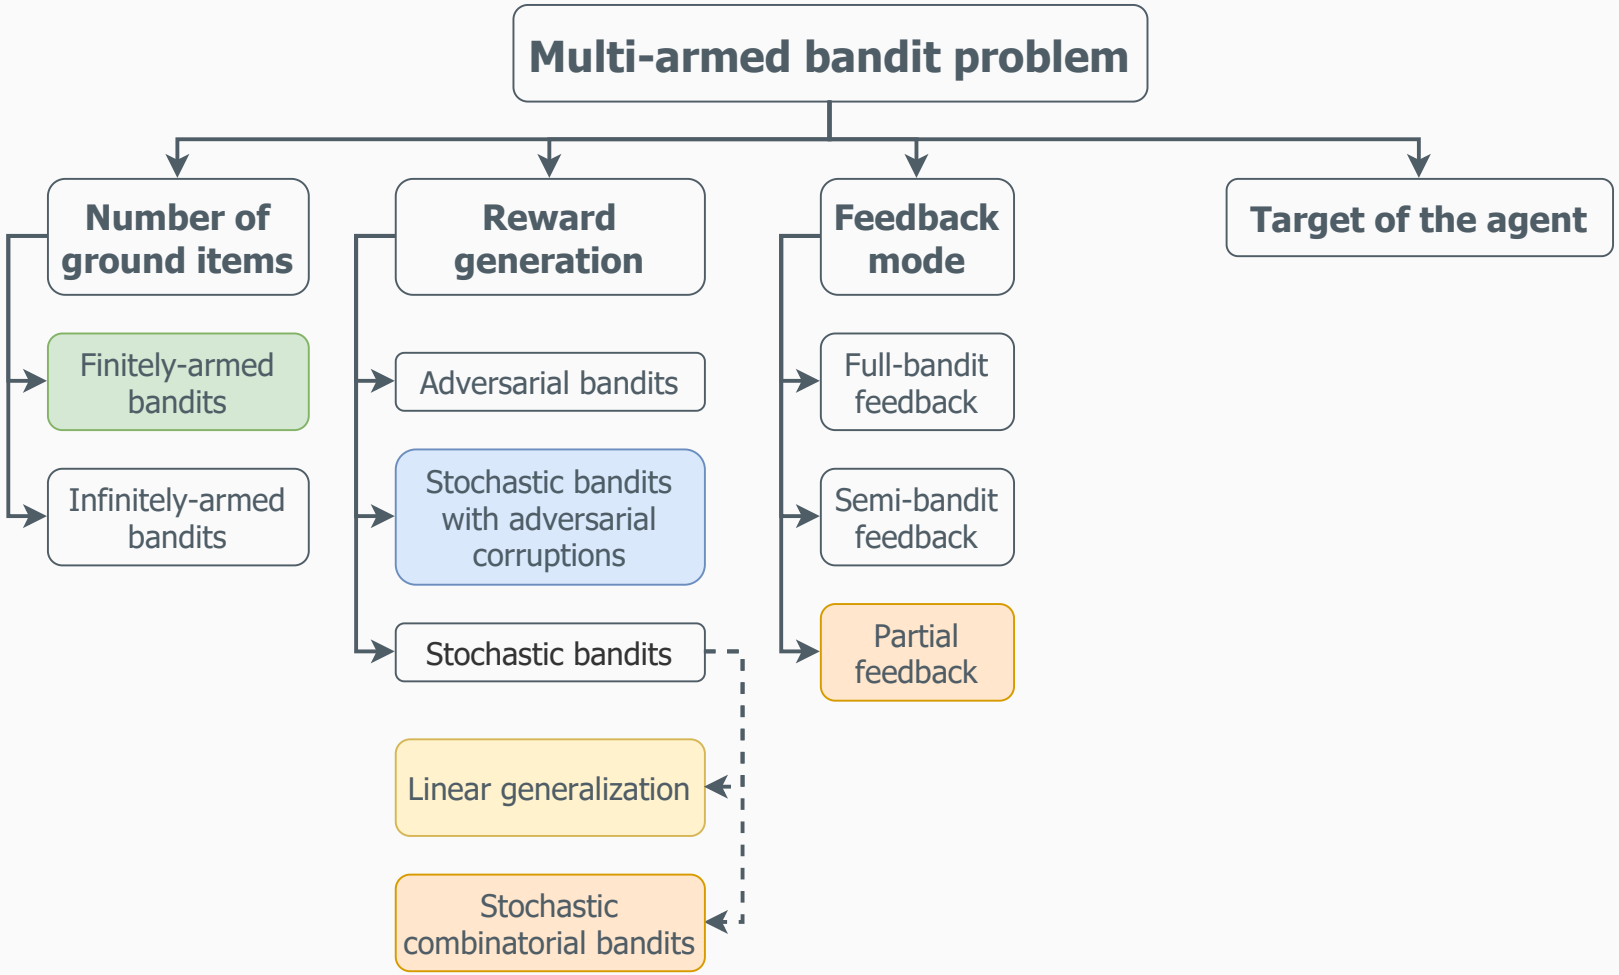

Classification

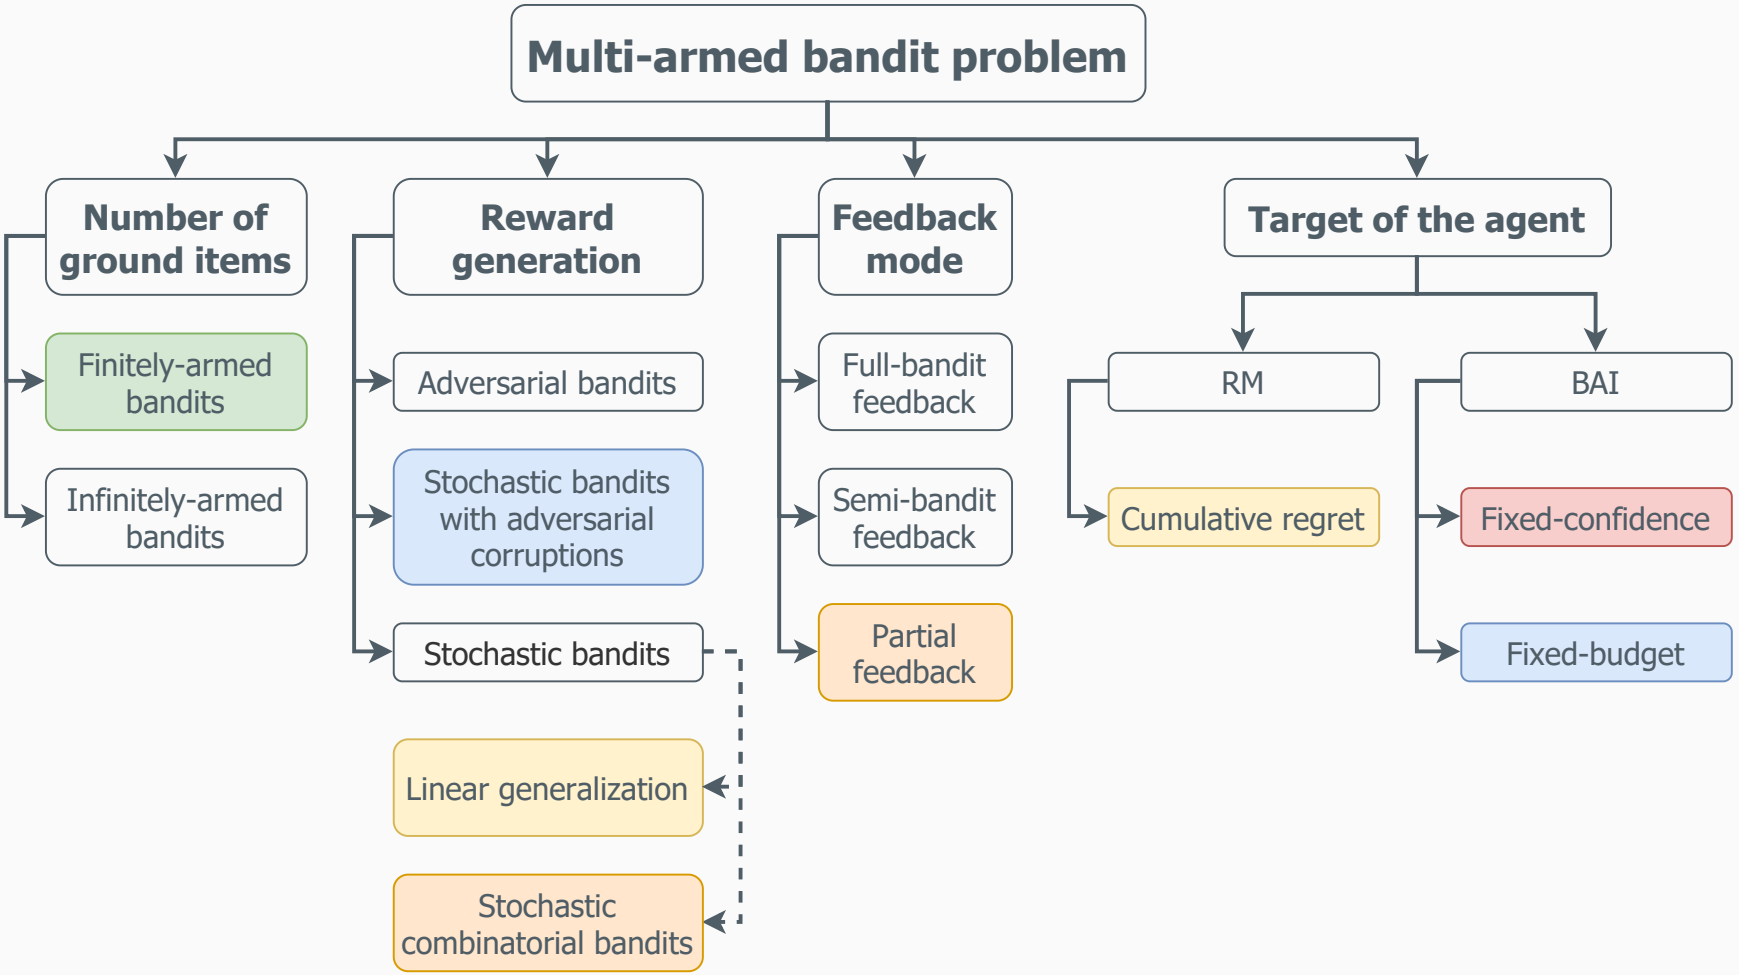

Formulation of MAB models

- **Ground set** — $\mathcal{S}$ consists of available arms.
- **Dynamics** — At each time step $t = 1, 2, \ldots$
  1. **Reward** $W_t(i)$ is associated with arm i.
  2. Agent **pulls** arm A_t
  3. Agent observes the corresponding **feedback** $O_t = f(\{W_t(i) : i \in A_t\})$.

Formulation of MAB models

♠ **Ground set** — $\mathcal{S}$ consists of available arms.

♠ **Dynamics** — At each time step $t = 1, 2, \ldots$

1. **Reward** $W_t(i)$ is associated with arm i.
2. Agent **pulls** arm A_t
3. Agent observes the corresponding **feedback** $O_t = f(\{W_t(i) : i \in A_t\})$.

♠ **Number of arms**

- Finite-armed bandits (Audibert et al., 2009; Agrawal and Goyal, 2012)

  Ground set $\mathcal{S}$ of L arms is indexed by $[L] = \{1, 2, \ldots, L\}$.
- Infinite-armed bandits (Berry et al., 1997)

  Related to the topic of Bayesian optimization

Stochastic bandits

- Each arm $i \in [L]$ is associated with an unknown distribution $\nu(i)$, mean $w(i)$, and variance $\sigma^2(i)$.
- $\{W_t(i)\}_{t=1}^T$ is the i.i.d. sequence of rewards associated with arm i during the T time steps.

Stochastic bandits

- Each arm $i \in [L]$ is associated with an unknown distribution $\nu(i)$, mean $w(i)$, and variance $\sigma^2(i)$.
- $\{W_t(i)\}_{t=1}^T$ is the i.i.d. sequence of rewards associated with arm i during the T time steps.

♠ Linear generalization (Abe and Long, 1999)

- $w(i) = x(i)^\top \beta$
- Feature vector $x(i) \in \mathbb{R}^d$ is known for each arm i, latent vector $\beta \in \mathbb{R}^d$ is not known.
- Reduces to standard bandits when $x(i) = e_i$, standard basis.

Stochastic bandits

- Each arm $i \in [L]$ is associated with an unknown distribution $\nu(i)$, mean $w(i)$, and variance $\sigma^2(i)$.
- $\{W_t(i)\}_{t=1}^T$ is the i.i.d. sequence of rewards associated with arm i during the T time steps.

♠ Linear generalization (Abe and Long, 1999)

- $w(i) = x(i)^\top \beta$
- Feature vector $x(i) \in \mathbb{R}^d$ is known for each arm i, latent vector $\beta \in \mathbb{R}^d$ is not known.
- Reduces to standard bandits when $x(i) = e_i$, standard basis.

♠ Stochastic combinatorial bandits

- Stardard setting: $|A_t| = 1$.
- Combinatorial setting: $|A_t| \geq 1$.

Feedback mode

- Full-bandit feedback
- Semi-bandit feedback
- Partial feedback

Feedback mode

♠ Full-bandit feedback

Agent only observes the sums of the realizations of all pulled arms (Rejwan and Mansour, 2020; Kuroki et al., 2020).

♠ Semi-bandit feedback

♠ Partial feedback

Feedback mode

♠ Full-bandit feedback

Agent only observes the sums of the realizations of all pulled arms (Rejwan and Mansour, 2020; Kuroki et al., 2020).

♠ Semi-bandit feedback

Agent observes realizations of all pulled arms (Mannor and Tsitsiklis, 2004; Kalyanakrishnan et al., 2012).

♠ Partial feedback

Feedback mode

♠ Full-bandit feedback

Agent only observes the sums of the realizations of all pulled arms (Rejwan and Mansour, 2020; Kuroki et al., 2020).

♠ Semi-bandit feedback

Agent observes realizations of all pulled arms (Mannor and Tsitsiklis, 2004; Kalyanakrishnan et al., 2012).

♠ Partial feedback

Agent only observes the realizations of a subset of pulled arms (Kveton et al., 2015b; Li et al., 2016).

Reward generation

♠ Stochastic bandits

Reward generation

♠ Stochastic bandits

♠ Stochastic bandits with adversarial corruptions

(Shen, 2019; Jun et al., 2018)

At each time step $t = 1, \ldots, T$:

1. Stochastic reward $W_t(i) \in [0, 1]$ is i.i.d. drawn for each arm i.

Reward generation

♠ Stochastic bandits

♠ Stochastic bandits with adversarial corruptions

(Shen, 2019; Jun et al., 2018)

At each time step $t = 1, \ldots, T$:

1. Stochastic reward $W_t(i) \in [0, 1]$ is i.i.d. drawn for each arm i.
2. Agent pulls arm i_t.
3. Adversary observes $\{W_t(i)\}_{i \in [L]}$ as well as i_t, and corrupts $W_t(i_t)$ with c_t:
$$\tilde{W}_t(i_t) = W_t(i_t) + c_t \in [0, 1].$$
but the norm of $\{c_t\}_{t=1}^T$ is suitably constrained.
4. Agent observes the corrupted reward $\tilde{W}_t(i_t)$.

Reward generation

♠ Stochastic bandits

♠ Stochastic bandits with adversarial corruptions

(Shen, 2019; Jun et al., 2018)

At each time step $t = 1, \ldots, T$:

1. Stochastic reward $W_t(i) \in [0, 1]$ is i.i.d. drawn for each arm i.
2. Agent pulls arm i_t.
3. Adversary observes $\{W_t(i)\}_{i \in [L]}$ as well as i_t, and corrupts $W_t(i_t)$ with c_t:
$$\tilde{W}_t(i_t) = W_t(i_t) + c_t \in [0, 1].$$
but the norm of $\{c_t\}_{t=1}^T$ is suitably constrained.
4. Agent observes the corrupted reward $\tilde{W}_t(i_t)$.

Reward generation

♠ Stochastic bandits

♠ Stochastic bandits with adversarial corruptions

(Shen, 2019; Jun et al., 2018)

At each time step $t = 1, \ldots, T$:

1. Stochastic reward $W_t(i) \in [0, 1]$ is i.i.d. drawn for each arm i.
2. Agent pulls arm i_t.
3. Adversary observes $\{W_t(i)\}_{i \in [L]}$ as well as i_t, and corrupts $W_t(i_t)$ with c_t:
$$\tilde{W}_t(i_t) = W_t(i_t) + c_t \in [0, 1].$$
but the norm of $\{c_t\}_{t=1}^T$ is suitably constrained.
4. Agent observes the corrupted reward $\tilde{W}_t(i_t)$.

♠ Adversarial/Non-stochastic bandits

(Auer et al., 2002b; Cesa-Bianchi and Lugosi, 2006)

- Rewards $\{W_t(i)\}_{t=1}$ of each arm i are not necessarily drawn independently from the same distribution.

♠ Stochastic bandits

♠ Stochastic bandits with adversarial corruptions

(Shen, 2019; Jun et al., 2018)

At each time step $t = 1, \ldots, T$:

1. Stochastic reward $W_t(i) \in [0, 1]$ is i.i.d. drawn for each arm i.
2. Agent pulls arm i_t.
3. Adversary observes $\{W_t(i)\}_{i \in [L]}$ as well as i_t, and corrupts $W_t(i_t)$ with c_t:
$$\tilde{W}_t(i_t) = W_t(i_t) + c_t \in [0, 1].$$
but the norm of $\{c_t\}_{t=1}^T$ is suitably constrained.
4. Agent observes the corrupted reward $\tilde{W}_t(i_t)$.

♠ Adversarial/Non-stochastic bandits

(Auer et al., 2002b; Cesa-Bianchi and Lugosi, 2006)

- Rewards $\{W_t(i)\}_{t=1}$ of each arm i are not necessarily drawn independently from the same distribution.

Stochastically constrained adversarial bandits (Zimmert and Seldin, 2021)

- $W_t(i)$ is a r.v. with mean $w_t(i)$, and gaps $\Delta_{i,j} = W_t(i) - W_t(j)$ are fixed.

Target of the agent

♠ Cumulative regret minimization

♠ Simple regret minimization

♠ Pure exploration/best arm identification (BAI)

Fixed-confidence setting

Fixed-budget setting

Target of the agent

♠ Cumulative regret minimization

Maximize the **cumulative** reward, i.e.,minimize the regret (the gap between the maximum cumulative reward and the reward obtained by the agent) (Agrawal and Goyal, 2012; Russo and Van Roy, 2014; Lai, 1987).

♠ Simple regret minimization

♠ Pure exploration/best arm identification (BAI)

Fixed-confidence setting

Fixed-budget setting

Target of the agent

♠ Cumulative regret minimization

Maximize the **cumulative** reward, i.e.,minimize the regret (the gap between the maximum cumulative reward and the reward obtained by the agent) (Agrawal and Goyal, 2012; Russo and Van Roy, 2014; Lai, 1987).

♠ Simple regret minimization

Maximize the **mean reward of the chosen arm** by the end of a fixed time horizon T (Carpentier and Valko, 2015).

♠ Pure exploration/best arm identification (BAI)

Fixed-confidence setting

Fixed-budget setting

Target of the agent

♠ Cumulative regret minimization

Maximize the **cumulative** reward, i.e.,minimize the regret (the gap between the maximum cumulative reward and the reward obtained by the agent) (Agrawal and Goyal, 2012; Russo and Van Roy, 2014; Lai, 1987).

♠ Simple regret minimization

Maximize the **mean reward of the chosen arm** by the end of a fixed time horizon T (Carpentier and Valko, 2015).

♠ Pure exploration/best arm identification (BAI)

Fixed-confidence setting Given a risk parameter δ, the agent aims to identify the best arm with probability $1-\delta$ in **minimal time steps** (Jamieson and Nowak, 2014; Kalyanakrishnan et al., 2012).

Fixed-budget setting

Target of the agent

♠ Cumulative regret minimization

Maximize the **cumulative** reward, i.e.,minimize the regret (the gap between the maximum cumulative reward and the reward obtained by the agent) (Agrawal and Goyal, 2012; Russo and Van Roy, 2014; Lai, 1987).

♠ Simple regret minimization

Maximize the **mean reward of the chosen arm** by the end of a fixed time horizon T (Carpentier and Valko, 2015).

♠ Pure exploration/best arm identification (BAI)

Fixed-confidence setting Given a risk parameter δ, the agent aims to identify the best arm with probability $1-\delta$ in **minimal time steps** (Jamieson and Nowak, 2014; Kalyanakrishnan et al., 2012).

Fixed-budget setting Given a budget constraint T, the agent aims to **maximize the confidence** of the chosen arm by the end of a fixed time horizon T (Auer et al., 2002a; Audibert and Bubeck, 2010; Carpentier and Locatelli, 2016).

Outline

1 **What is multi-armed bandits (MAB)?**

- Example — Cascading bandits

Example — Cascading bandits (Kveton et al., 2015a)

♠ Online recommender system

- seek to select a small list of items to the user over time.

Example — Cascading bandits (Kveton et al., 2015a)

♠ Online recommender system

- seek to select a small list of items to the user over time.

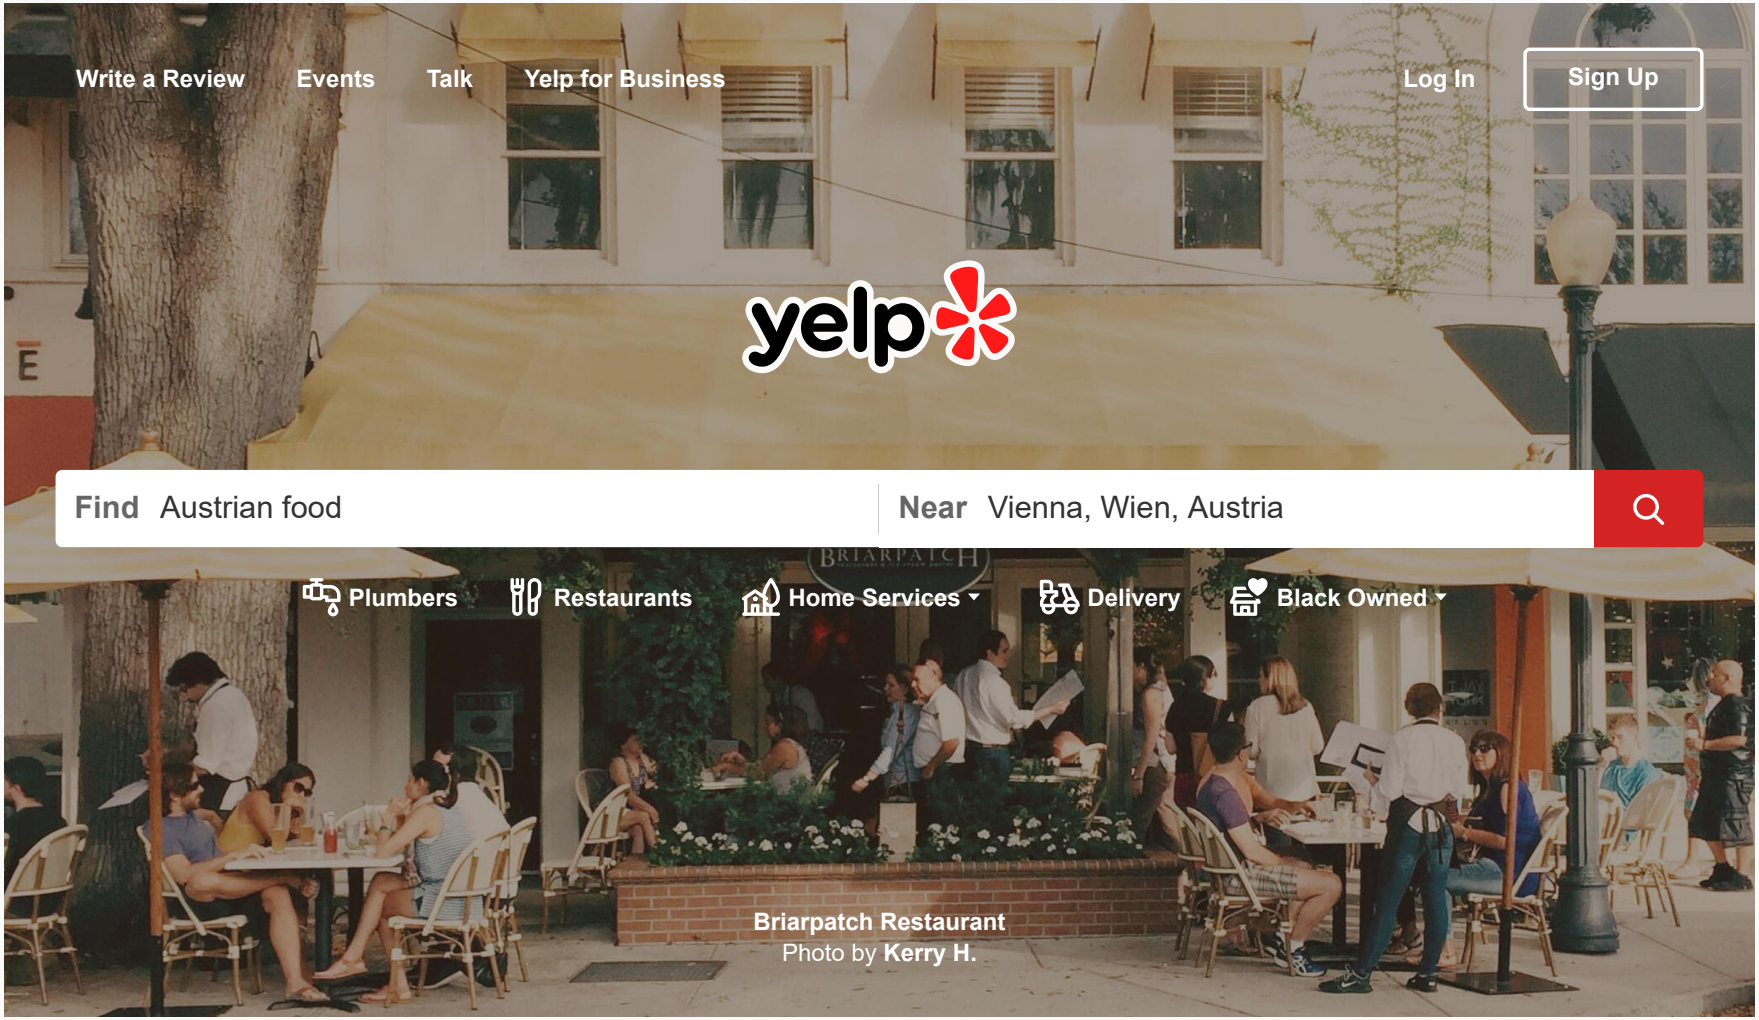

Example — Cascading bandits (Kveton et al., 2015a)

yelp

Austrian food | Vienna, Wien, Austria

For Businesses | Write a Review | Log In | Sign Up

Restaurants | Home Services | Auto Services | More

Filters

€ | €€ | €€€ | €€€€

Suggested

- [ ] Open Now 4:04 PM

Category

Austrian | Bars | Cafes | Gastropubs

See all

Features

- [ ] Good for Groups
- [ ] Takes Reservations
- [ ] Outdoor Seating
- [ ] Good for Kids

See all

Neighborhoods

- [ ] Floridsdorf
- [ ] Innere Stadt
- [ ] Leopoldstadt
- [ ] Landstraße

See all

Wien > Restaurants > Austrian food

Best Austrian food in Vienna, Wien, Austria

Sort: **Recommended**

All | Price | Open Now

1. Gasthaus Pöschl

236

Gastropubs | Austrian | €€ • Innere Stadt

Closed until Noon

"Really nice service and traditional **Austrian** food. The salads are generous & the goulash delicious" **more**

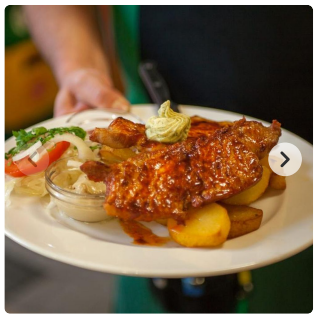

2. Gasthaus Kopp

68

Austrian | Beisl | € • Brigittenau

Open until Midnight

"Take the U Bahn tu the S Bahn and walk 4 blocks to get to **Austrian** food heaven. I was alone but" **more**

Example — Cascading bandits (Kveton et al., 2015a)

♠ Online recommender system

- Seek to select a small list of items to the user over time.
- How to maximize the 'reward' over several rounds of recommendation?

  — Regret Minimization (RM)

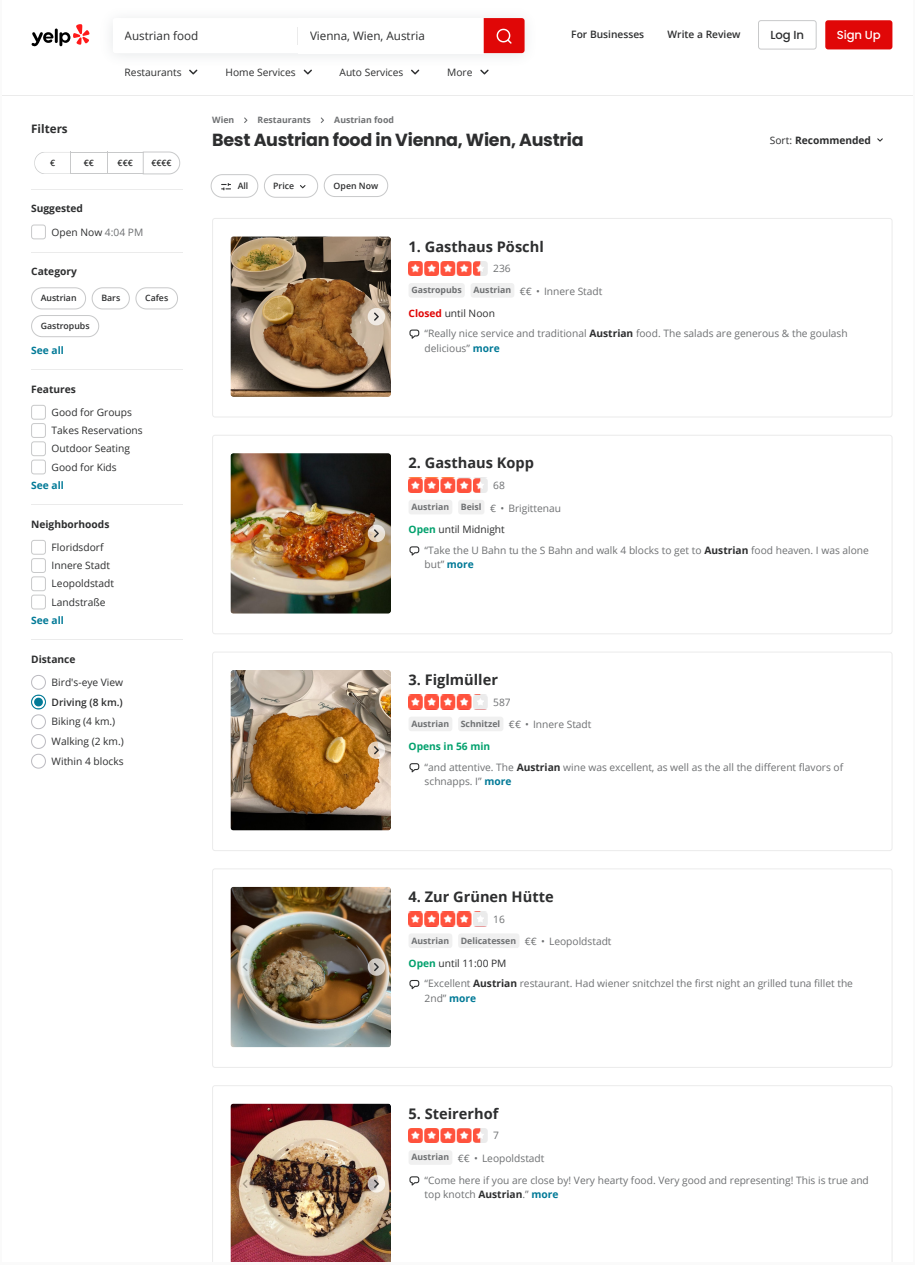

Example — Cascading bandits (Kveton et al., 2015a)

♠ Online recommender system

- Seek to select a small list of items to the user over time.
- How to maximize the 'reward' over several rounds of recommendation?
  — Regret Minimization (RM)
- How to select an attractive list of items after several rounds of recommendation?
  — Pure Exploration/
  Best Arm Identification (BAI)

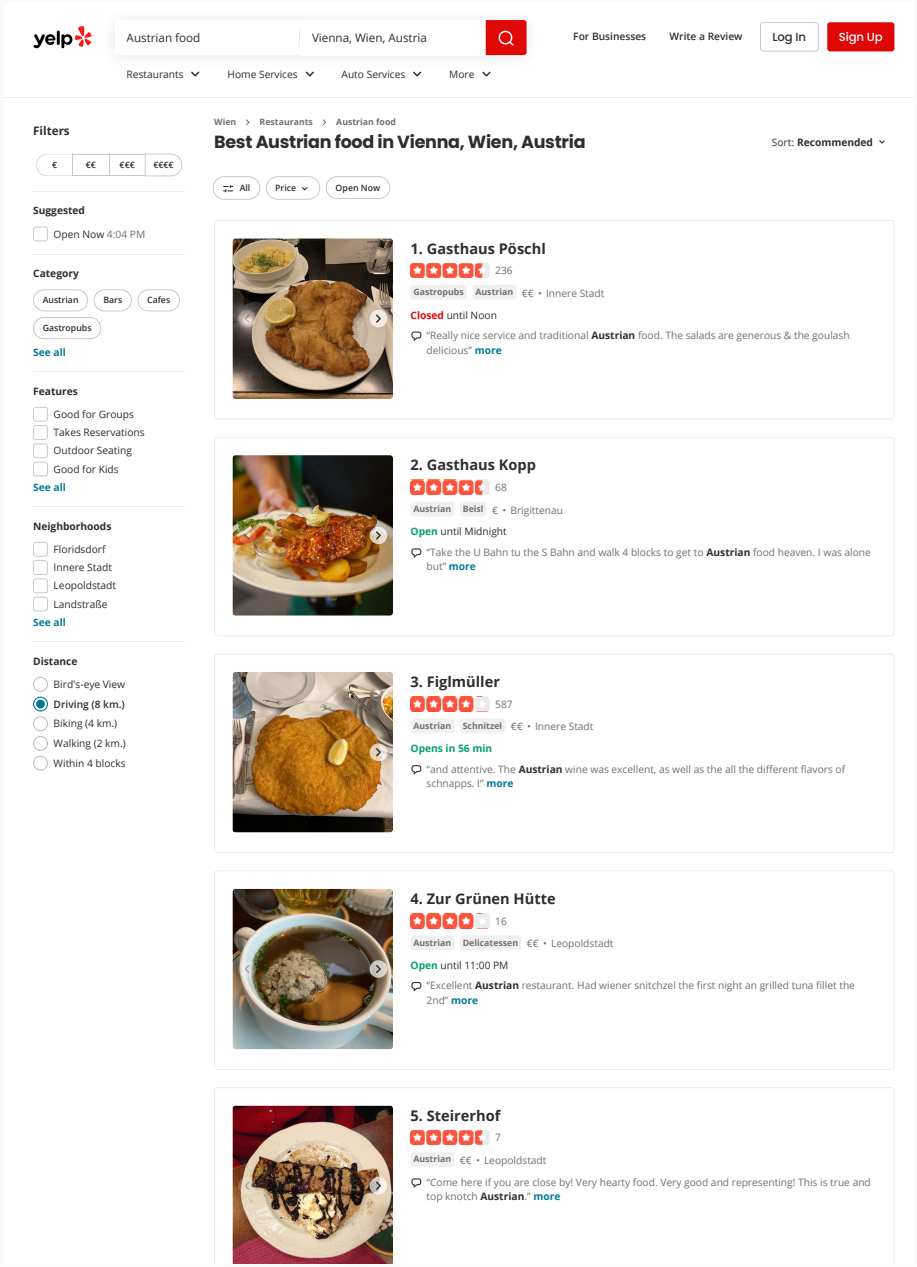

Example — Cascading bandits (Kveton et al., 2015a)

Ground set

A finite set of all available arms $[L] := \{1, \ldots, L\}$.

Click probability/weight of item $i \in [L]$

Arm i attracts the user with probability $w(i) \in [0, 1]$.

Example — Cascading bandits (Kveton et al., 2015a)

Ground set

A finite set of all available arms $[L] := \{1, \ldots, L\}$.

Click probability/weight of item $i \in [L]$

Arm i attracts the user with probability $w(i) \in [0, 1]$.

- Standard setting: $\boldsymbol{w} := \{w(i)\}_{i=1}^{L}$ are not known.
- Linear generalization: $w(i) = x(i)^\top \beta$
  Feature vector $x(i)$ is known for each arm i, latent vector $\beta \in \mathbb{R}^d$ is not known.

Example — Cascading bandits (Kveton et al., 2015a)

Ground set

A finite set of all available arms $[L] := \{1, \ldots, L\}$.

Click probability/weight of item $i \in [L]$

Arm i attracts the user with probability $w(i) \in [0, 1]$.

- Standard setting: $\boldsymbol{w} := \{w(i)\}_{i=1}^{L}$ are not known.
- Linear generalization: $w(i) = x(i)^\top \beta$
  Feature vector $x(i)$ is known for each arm i, latent vector $\beta \in \mathbb{R}^d$ is not known.

Whether arm i is clicked at time t

This is revealed by a random variable $W_t(i) \sim \mathrm{Bern}\,(w(i))$.

- $W_t(i) = 1$ iff the user observes and clicks on i at time t.
- $W_t(i) = 0$ iff the user observes but does not click on i at time t.

Example — Cascading bandits (Kveton et al., 2015a)

Ground set

A finite set of all available arms $[L] := \{1, \ldots, L\}$.

Click probability/weight of item $i \in [L]$

Arm i attracts the user with probability $w(i) \in [0, 1]$.

- Standard setting: $\boldsymbol{w} := \{w(i)\}_{i=1}^{L}$ are not known.
- Linear generalization: $w(i) = x(i)^\top \beta$
  Feature vector $x(i)$ is known for each arm i, latent vector $\beta \in \mathbb{R}^d$ is not known.

Whether arm i is clicked at time t

This is revealed by a random variable $W_t(i) \sim \mathrm{Bern}\,(w(i))$.

- $W_t(i) = 1$ iff the user observes and clicks on i at time t.
- $W_t(i) = 0$ iff the user observes but does not click on i at time t.
- ♦ $W_t(i)$'s are only observed for some arms.

Example — Cascading bandits (Kveton et al., 2015a)

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

Example — Cascading bandits (Kveton et al., 2015a)

$$L = 9$$

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

Example — Cascading bandits (Kveton et al., 2015a)

Recommendation 

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

Example — Cascading bandits (Kveton et al., 2015a)

Recommendation 

$$K = 5$$

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

Example — Cascading bandits (Kveton et al., 2015a)

Recommendation

Attractiveness
$W_t(i)$

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

2. The user examines the arms from i_1^t to i_K^t:
   - If she is attracted by an item, clicks on it;
   - If not, she skips to the next item and checks if it is attractive;
   - Process stops when she clicks on one item or when she comes to the end of the list.

Example — Cascading bandits (Kveton et al., 2015a)

Recommendation		
Attractiveness	$\times$	
$W_t(i)$	0	

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

2. The user examines the arms from i_1^t to i_K^t:
   - If she is attracted by an item, clicks on it;
   - If not, she skips to the next item and checks if it is attractive;
   - Process stops when she clicks on one item or when she comes to the end of the list.

Example — Cascading bandits (Kveton et al., 2015a)

Recommendation			
Attractiveness	×	×	
$W_t(i)$	0	0	

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

2. The user examines the arms from i_1^t to i_K^t:
   - If she is attracted by an item, clicks on it;
   - If not, she skips to the next item and checks if it is attractive;
   - Process stops when she clicks on one item or when she comes to the end of the list.

Example — Cascading bandits (Kveton et al., 2015a)

			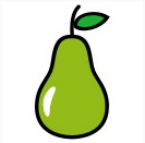	
Recommendation				
Attractiveness	$\times$	$\times$	$\times$	
$W_t(i)$	0	0	0	

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

2. The user examines the arms from i_1^t to i_K^t:
   - If she is attracted by an item, clicks on it;
   - If not, she skips to the next item and checks if it is attractive;
   - Process stops when she clicks on one item or when she comes to the end of the list.

Example — Cascading bandits (Kveton et al., 2015a)

Recommendation				
Attractiveness	$\times$	$\times$	$\times$	$\surd$
$W_t(i)$	0	0	0	1

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

2. The user examines the arms from i_1^t to i_K^t:
   - If she is attracted by an item, clicks on it;
   - If not, she skips to the next item and checks if it is attractive;
   - Process stops when she clicks on one item or when she comes to the end of the list.

Example — Cascading bandits (Kveton et al., 2015a)

Recommendation					
Attractiveness	$\times$	$\times$	$\times$	$\surd$	$?$
$W_t(i)$	0	0	0	1	$?$

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

2. The user examines the arms from i_1^t to i_K^t:
   - If she is attracted by an item, clicks on it;
   - If not, she skips to the next item and checks if it is attractive;
   - Process stops when she clicks on one item or when she comes to the end of the list.

Example — Cascading bandits (Kveton et al., 2015a)

Recommendation					
Attractiveness	$\times$	$\times$	$\times$	$\surd$	?
$W_t(i)$	0	0	0	1	?

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

2. The user examines the arms from i_1^t to i_K^t:
   - If she is attracted by an item, clicks on it;
   - If not, she skips to the next item and checks if it is attractive;
   - Process stops when she clicks on one item or when she comes to the end of the list.

♠ **Combinatorial** bandits ♡ **Partial** feedback

Example — Cascading bandits (Kveton et al., 2015a)

Recommendation					
Attractiveness	$\times$	$\times$	$\times$	$\surd$	?
$W_t(i)$	0	0	0	1	?

For each time step $t = 1, 2, \ldots$

1. The agent selects a list of K arms $S_t := (i_1^t, \ldots, i_K^t) \in [L]^{(K)}$ to the user, where $[L]^{(K)} = \{\text{all } K\text{-permutations of } [L]\}$;

2. The user examines the arms from i_1^t to i_K^t:
   - If she is attracted by an item, clicks on it;
   - If not, she skips to the next item and checks if it is attractive;
   - Process stops when she clicks on one item or when she comes to the end of the list.

♠ **Combinatorial** bandits ♡ **Partial** feedback ♣ **Standard** setting & **Linear** generalization

Outline

1 What is multi-armed bandits (MAB)?

Pure exploration/BAI settings

Fixed-confidence setting

- Given a risk parameter δ, the agent aims to identify the best arm with probability $1-\delta$ in minimal time steps.
  (Jamieson and Nowak, 2014; Kalyanakrishnan et al., 2012)

Fixed-budget setting

- Given a budget constraint T, the agent aims to maximize the confidence of the chosen arm by the end of a fixed time horizon T.
  (Auer et al., 2002a; Audibert and Bubeck, 2010; Carpentier and Locatelli, 2016)

Pure exploration in stochastic bandits

- Ground set $\mathcal{S} = [L]$ consists of L available arms.
- Each arm $i \in [L]$ is associated with an unknown distribution $\nu(i)$, mean $w(i)$, and variance $\sigma^2(i)$.
- $\{W_t(i)\}_{t=1}^T$ is the i.i.d. sequence of rewards associated with arm i during the T time steps.

Pure exploration in stochastic bandits

- Ground set $\mathcal{S} = [L]$ consists of L available arms.
- Each arm $i \in [L]$ is associated with an unknown distribution $\nu(i)$, mean $w(i)$, and variance $\sigma^2(i)$.
- $\{W_t(i)\}_{t=1}^T$ is the i.i.d. sequence of rewards associated with arm i during the T time steps.

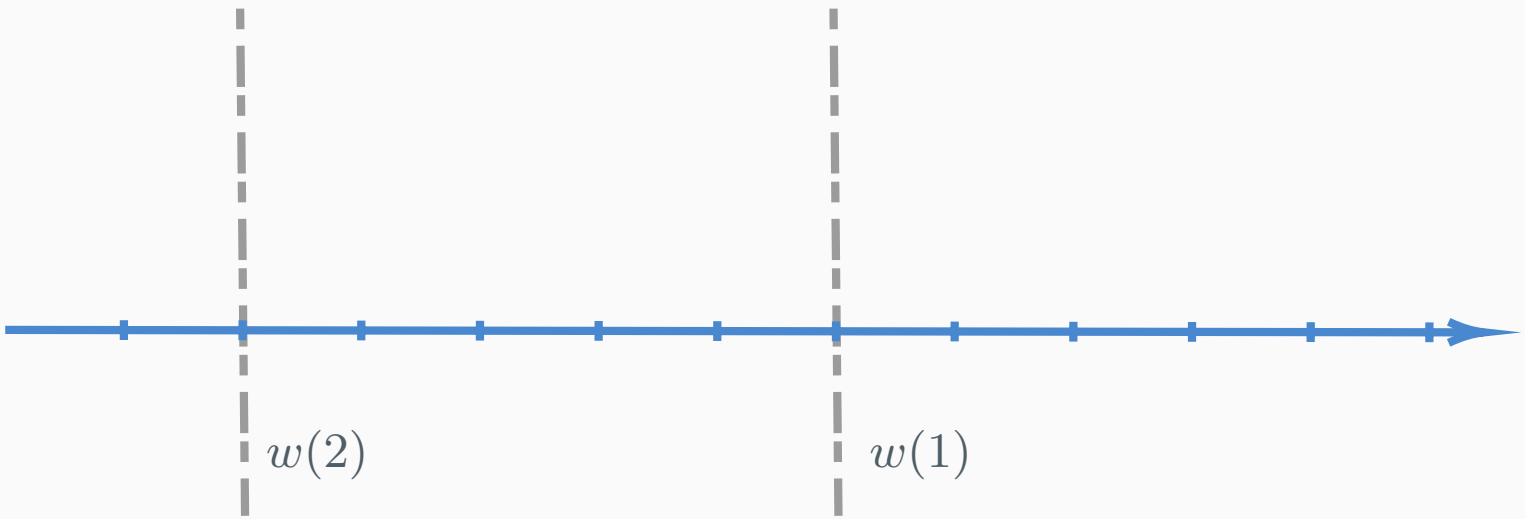

Pure exploration in stochastic bandits

- Ground set $\mathcal{S} = [L]$ consists of L available arms.
- Each arm $i \in [L]$ is associated with an unknown distribution $\nu(i)$, mean $w(i)$, and variance $\sigma^2(i)$.
- $\{W_t(i)\}_{t=1}^T$ is the i.i.d. sequence of rewards associated with arm i during the T time steps.

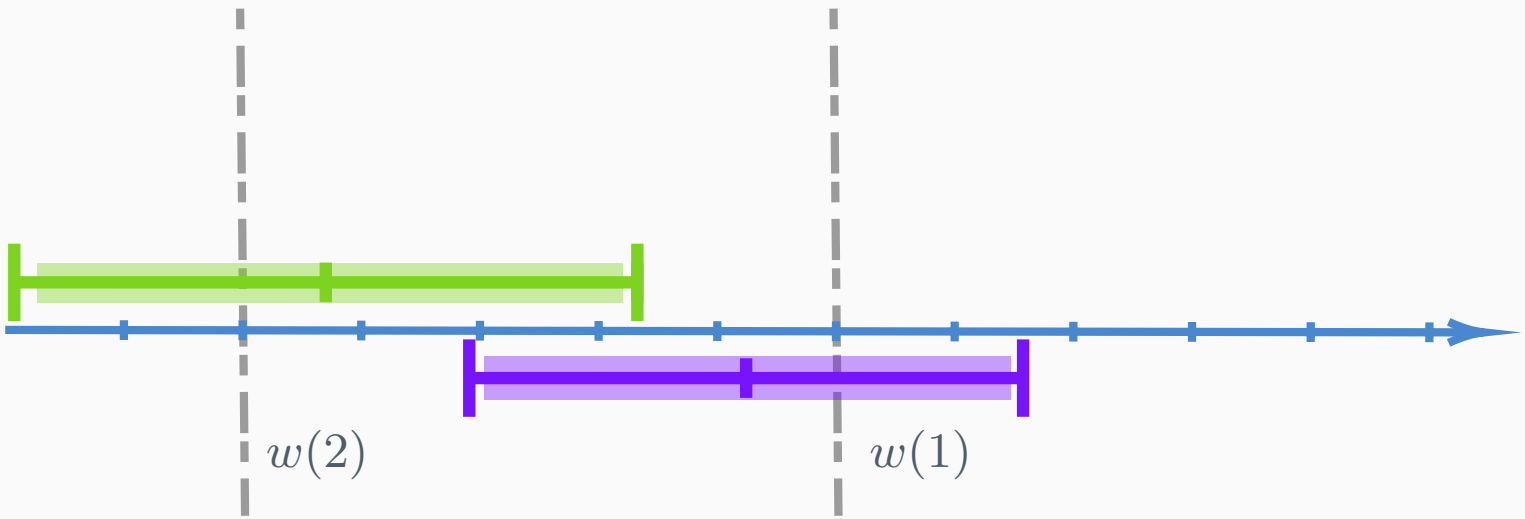

Pure exploration in stochastic bandits

- Ground set $\mathcal{S} = [L]$ consists of L available arms.
- Each arm $i \in [L]$ is associated with an unknown distribution $\nu(i)$, mean $w(i)$, and variance $\sigma^2(i)$.
- $\{W_t(i)\}_{t=1}^T$ is the i.i.d. sequence of rewards associated with arm i during the T time steps.

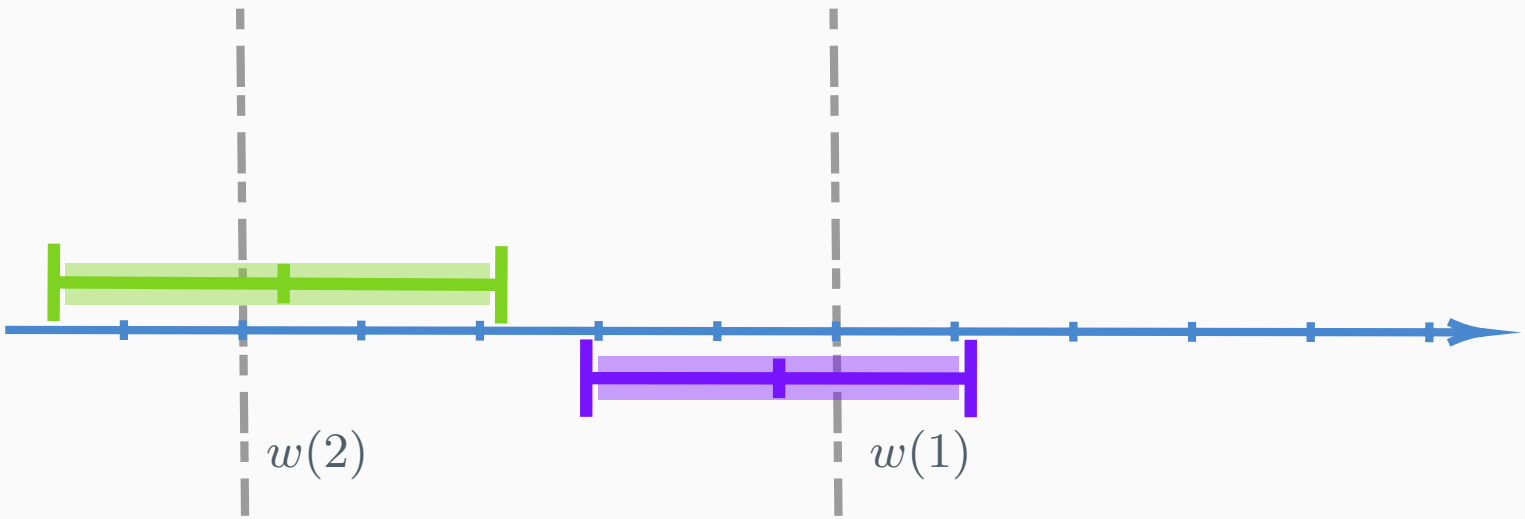

- Optimal arm

$$1 = i^* = \arg\max_{i \in [L]} w(i)$$

  - Without loss of generality, assume

$$w(1) > w(2) \geq w(3) \geq \ldots \geq w(L).$$

Pure exploration in stochastic bandits

- Optimal arm

$$1 = i^* = \arg\max_{i \in [L]} w(i)$$

  - Without loss of generality, assume

$$w(1) > w(2) \geq w(3) \geq \ldots \geq w(L).$$

- Gaps to optimality

$$\Delta_i = w(1) - w(i) \quad \forall\, i \neq 1, \quad \Delta_1 = \Delta_2.$$

Pure exploration in stochastic bandits

- Optimal arm

$$1 = i^* = \arg\max_{i \in [L]} w(i)$$

  - Without loss of generality, assume

$$w(1) > w(2) \geq w(3) \geq \ldots \geq w(L).$$

- Gaps to optimality

$$\Delta_i = w(1) - w(i) \quad \forall\, i \neq 1, \quad \Delta_1 = \Delta_2.$$

- Hardness parameters

$$H_1 = \sum_{i=1}^{L} \frac{1}{\Delta_i^2}, \quad H_2 = \max_{i \in [L]} \frac{i}{\Delta_i^2}.$$

Concentration inequalities: can we estimate $w(i)$ well?

Theorem 2.1 (Standard multiplicative variant of the Chernoff-Hoeffding bound; Dubhashi and Panconesi (2009), Theorem 1.1)

Suppose that $X_1, \dots, X_T$ are independent $[0,1]$-valued random variables, and let $X = \sum_{t=1}^{T} X_t$. Then for any $\varepsilon \in (0,1)$,

$$\Pr(X - \mathbb{E}[X] \geq \varepsilon \mathbb{E}[X]) \leq \exp\left(-\frac{\varepsilon^2}{3}\mathbb{E}X\right),$$

$$\Pr(X - \mathbb{E}[X] \leq -\varepsilon \mathbb{E}[X]) \leq \exp\left(-\frac{\varepsilon^2}{3}\mathbb{E}X\right).$$

BAI algorithm

A deterministic and non-anticipatory online algorithm consists in a triple $\pi := ((\pi_t)_t, \mathcal{T}^\pi, \phi^\pi)$

- sampling rule $(\pi_t)_t$: which arm S_t^π to pull at time step t
- stopping rule $\mathcal{T}^\pi$: when to stop
- recommendation rule ϕ^π: which arm $\hat{S}^\pi$ to choose eventually

BAI algorithm

A deterministic and non-anticipatory online algorithm consists in a triple $\pi := ((\pi_t)_t, \mathcal{T}^\pi, \phi^\pi)$

- sampling rule $(\pi_t)_t$: which arm S_t^π to pull at time step t
  S_t^π is $\mathcal{F}_{t-1}$-measurable, observation history $\mathcal{F}_t := \sigma(S_1^\pi, \boldsymbol{O}_1^\pi, \ldots, S_t^\pi, \boldsymbol{O}_t^\pi)$;
- stopping rule $\mathcal{T}^\pi$: when to stop
- recommendation rule ϕ^π: which arm $\hat{S}^\pi$ to choose eventually

BAI algorithm

A deterministic and non-anticipatory online algorithm consists in a triple $\pi := ((\pi_t)_t, \mathcal{T}^\pi, \phi^\pi)$

- sampling rule $(\pi_t)_t$: which arm S_t^π to pull at time step t
  S_t^π is $\mathcal{F}_{t-1}$-measurable, observation history $\mathcal{F}_t := \sigma(S_1^\pi, \boldsymbol{O}_1^\pi, \ldots, S_t^\pi, \boldsymbol{O}_t^\pi)$;
- stopping rule $\mathcal{T}^\pi$: when to stop
  stopping time $\mathcal{T}^\pi$ with respect to $(\mathcal{F}_t)_{t\in\mathbb{N}}$ satisfying $\mathbb{P}(\mathcal{T}^\pi < \infty) = 1$;
- recommendation rule ϕ^π: which arm $\hat{S}^\pi$ to choose eventually

BAI algorithm

A deterministic and non-anticipatory online algorithm consists in a triple $\pi := ((\pi_t)_t, \mathcal{T}^\pi, \phi^\pi)$

- sampling rule $(\pi_t)_t$: which arm S_t^π to pull at time step t
  S_t^π is $\mathcal{F}_{t-1}$-measurable, observation history $\mathcal{F}_t := \sigma(S_1^\pi, \boldsymbol{O}_1^\pi, \ldots, S_t^\pi, \boldsymbol{O}_t^\pi)$;
- stopping rule $\mathcal{T}^\pi$: when to stop
  stopping time $\mathcal{T}^\pi$ with respect to $(\mathcal{F}_t)_{t\in\mathbb{N}}$ satisfying $\mathbb{P}(\mathcal{T}^\pi < \infty) = 1$;
- recommendation rule ϕ^π: which arm $\hat{S}^\pi$ to choose eventually
  $\mathcal{F}_{\mathcal{T}^\pi}$-measurable.

BAI algorithm

A deterministic and non-anticipatory online algorithm consists in a triple $\pi := ((\pi_t)_t, \mathcal{T}^\pi, \phi^\pi)$

- sampling rule $(\pi_t)_t$: which arm S_t^π to pull at time step t
  S_t^π is $\mathcal{F}_{t-1}$-measurable, observation history $\mathcal{F}_t := \sigma(S_1^\pi, \boldsymbol{O}_1^\pi, \ldots, S_t^\pi, \boldsymbol{O}_t^\pi)$;
- stopping rule $\mathcal{T}^\pi$: when to stop
  stopping time $\mathcal{T}^\pi$ with respect to $(\mathcal{F}_t)_{t\in\mathbb{N}}$ satisfying $\mathbb{P}(\mathcal{T}^\pi < \infty) = 1$;
- recommendation rule ϕ^π: which arm $\hat{S}^\pi$ to choose eventually
  $\mathcal{F}_{\mathcal{T}^\pi}$-measurable.

$\mathcal{T}^\pi$

- Fixed-confidence setting: Time complexity of π (to minimize).
- Fixed-budget setting: $\mathcal{T}^\pi = \mathcal{T}$ (fixed).

Outline

1 What is multi-armed bandits (MAB)?

BAI: fixed-confidence

- δ-PAC algorithm: find the optimal arm with probability at least $1-\delta$

BAI: fixed-confidence

- δ-PAC algorithm: find the optimal arm with probability at least $1-\delta$

Theoretical study

- ▲ Propose a δ-PAC algorithm and upper bound its time complexity
- ▼ Derive a lower bound on the time complexity of any δ-PAC algorithm
- Evaluate theoretical findings with experiments

BAI: fixed-confidence

- δ-PAC algorithm: find the optimal arm with probability at least $1-\delta$

Theoretical study

- ▲ Propose a δ-PAC algorithm and upper bound its time complexity
- ▼ Derive a lower bound on the time complexity of any δ-PAC algorithm
- Evaluate theoretical findings with experiments

Simple pure exploration in stochastic bandits

- to identify the best arm with the largest mean:

$$i^* = \arg\max_{i \in [L]} w(i)$$

BAI: fixed-confidence

- δ-PAC algorithm: find the optimal arm with probability at least $1-\delta$

Theoretical study

- ▲ Propose a δ-PAC algorithm and upper bound its time complexity
- ▼ Derive a lower bound on the time complexity of any δ-PAC algorithm
- Evaluate theoretical findings with experiments

Simple pure exploration in stochastic bandits

- to identify the best arm with the largest mean:

$$i^* = \arg\max_{i \in [L]} w(i)$$

♠ Successive elimination

Successive Elimination, Median Elimination (Even-Dar et al., 2002)

BAI: fixed-confidence

- δ-PAC algorithm: find the optimal arm with probability at least $1-\delta$

Theoretical study

- ▲ Propose a δ-PAC algorithm and upper bound its time complexity
- ▼ Derive a lower bound on the time complexity of any δ-PAC algorithm
- Evaluate theoretical findings with experiments

Simple pure exploration in stochastic bandits

- to identify the best arm with the largest mean:

$$i^* = \arg\max_{i \in [L]} w(i)$$

♠ Successive elimination

Successive Elimination, Median Elimination (Even-Dar et al., 2002)

♠ Track optimal allocation

Track & Stop (Garivier and Kaufmann, 2016)

Algorithm 1: Successive Elimination(δ) (Even-Dar et al., 2002)

Algorithm 1: SUCCESSIVE ELIMINATION(δ) (Even-Dar et al., 2002)

1: Input: Set $t = 1$ and survival set $S = [L]$.
2: Let $\hat{w}_i^t$ be the average reward of arm i by time t.
3: Set $\hat{w}_i^1 = 0$ for all arm $i \in [L]$.

Algorithm 1: SUCCESSIVE ELIMINATION(δ) (Even-Dar et al., 2002)

1: Input: Set $t = 1$ and survival set $S = [L]$.
2: Let $\hat{w}_i^t$ be the average reward of arm i by time t.
3: Set $\hat{w}_i^1 = 0$ for all arm $i \in [L]$.
4: Sample each arm $i \in S$ once and update $\hat{w}_i^t$ (average reward of arm i).
5: Let $\hat{w}_{\max}^t = \max_{i \in [L]} \hat{w}_i^t$ and confidence radius $\alpha_t = \sqrt{\frac{\log(cLt^2/\delta)}{t}}$.
6: For each arm $i \in S$ such that $\hat{w}_{\max}^t - \hat{w}_i^t \geq 2\alpha_t$, set $S = S \setminus \{i\}$.

Algorithm 1: SUCCESSIVE ELIMINATION(δ) (Even-Dar et al., 2002)

1: Input: Set $t = 1$ and survival set $S = [L]$.
2: Let $\hat{w}_i^t$ be the average reward of arm i by time t.
3: Set $\hat{w}_i^1 = 0$ for all arm $i \in [L]$.
4: Sample each arm $i \in S$ once and update $\hat{w}_i^t$ (average reward of arm i).
5: Let $\hat{w}_{\max}^t = \max_{i \in [L]} \hat{w}_i^t$ and confidence radius $\alpha_t = \sqrt{\frac{\log(cLt^2/\delta)}{t}}$.
6: For each arm $i \in S$ such that $\hat{w}_{\max}^t - \hat{w}_i^t \geq 2\alpha_t$, set $S = S \setminus \{i\}$.
7: $t = t + 1$.
8: If $|S| > 1$, the Go to Step 4, Else output S.

Algorithm 1: SUCCESSIVE ELIMINATION(δ) (Even-Dar et al., 2002)

1: Input: Set $t = 1$ and survival set $S = [L]$.
2: Let $\hat{w}_i^t$ be the average reward of arm i by time t.
3: Set $\hat{w}_i^1 = 0$ for all arm $i \in [L]$.
4: Sample each arm $i \in S$ once and update $\hat{w}_i^t$ (average reward of arm i).
5: Let $\hat{w}_{\max}^t = \max_{i \in [L]} \hat{w}_i^t$ and confidence radius $\alpha_t = \sqrt{\frac{\log(cLt^2/\delta)}{t}}$.
6: For each arm $i \in S$ such that $\hat{w}_{\max}^t - \hat{w}_i^t \geq 2\alpha_t$, set $S = S \setminus \{i\}$.
7: $t = t + 1$.
8: If $|S| > 1$, the Go to Step 4, Else output S.

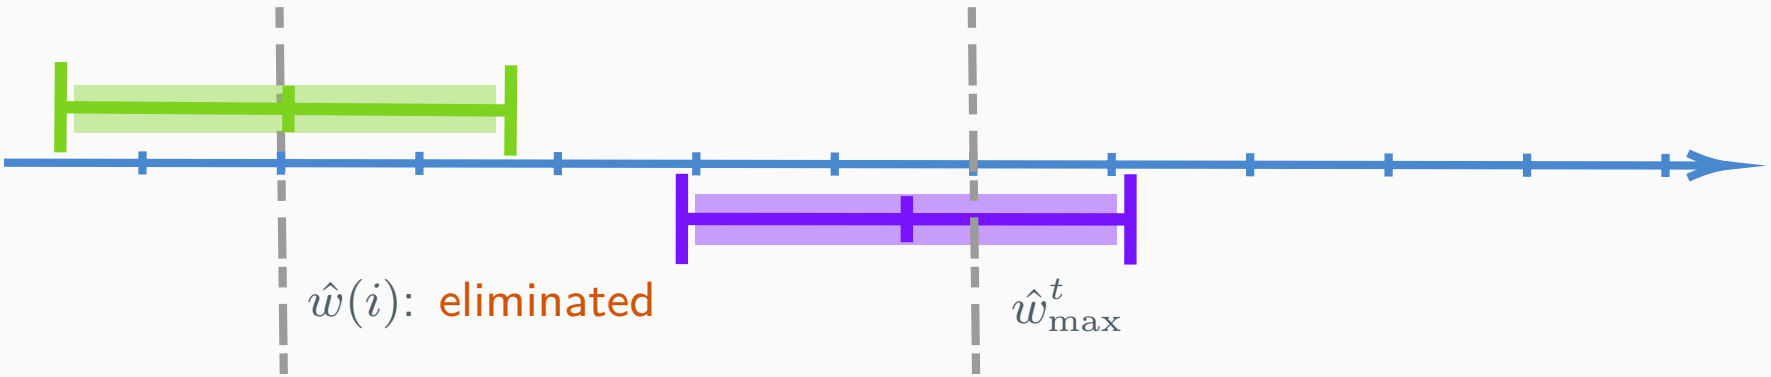

Algorithm 1: SUCCESSIVE ELIMINATION(δ) (Even-Dar et al., 2002)

1: Input: Set $t = 1$ and survival set $S = [L]$.
2: Let $\hat{w}_i^t$ be the average reward of arm i by time t.
3: Set $\hat{w}_i^1 = 0$ for all arm $i \in [L]$.
4: Sample each arm $i \in S$ once and update $\hat{w}_i^t$ (average reward of arm i).
5: Let $\hat{w}_{\max}^t = \max_{i \in [L]} \hat{w}_i^t$ and confidence radius $\alpha_t = \sqrt{\frac{\log(cLt^2/\delta)}{t}}$.
6: For each arm $i \in S$ such that $\hat{w}_{\max}^t - \hat{w}_i^t \geq 2\alpha_t$, set $S = S \setminus \{i\}$.
7: $t = t + 1$.
8: If $|S| > 1$, the Go to Step 4, Else output S.

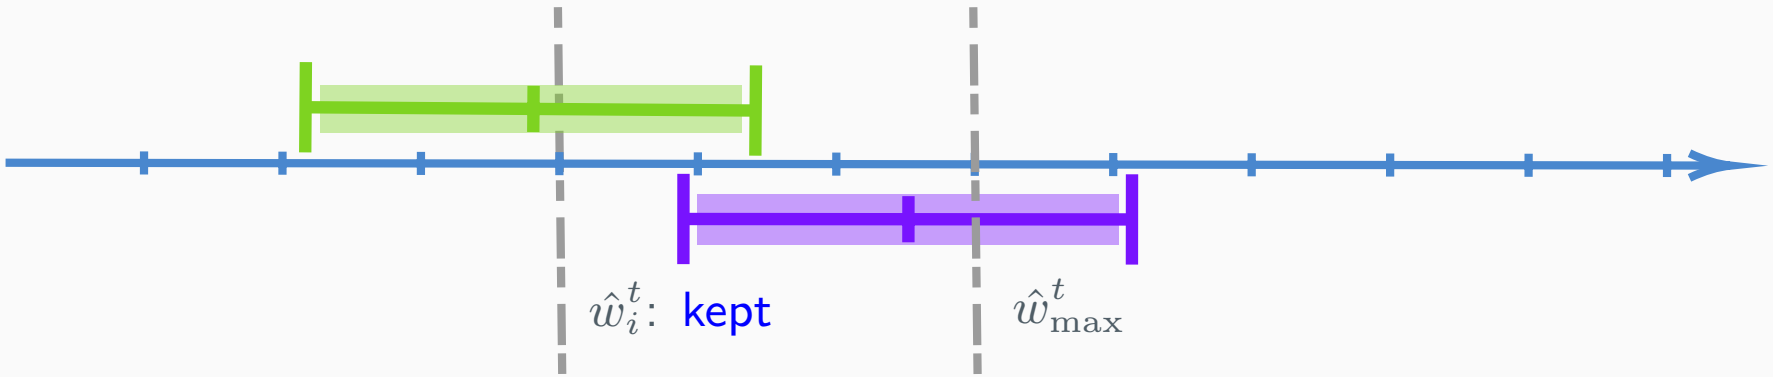

Step 1. Concentration inequality:

$$\Pr\left(\cup_{i\in[L]} \cup_{t\in\mathbb{N}} \left\{ |\hat{w}_i^t - w(i)| > \alpha_t \right\} \right) \leq \delta.$$

Step 1. Concentration inequality:

$$\Pr\left(\cup_{i\in[L]} \cup_{t\in\mathbb{N}} \left\{ |\hat{w}_i^t - w(i)| > \alpha_t \right\} \right) \leq \delta.$$

Step 2. Assume $\cap_{i\in[L]} \cap_{t\in\mathbb{N}} \left\{ |\hat{w}_i^t - w(i)| \leq \alpha_t \right\}$ holds. Recall that we eliminate arm i when $\hat{w}_{\max}^t - \hat{w}_i^t \geq 2\alpha_t$. Since

$$\hat{w}_{\max}^t - \hat{w}_i^t \geq \hat{w}_1^t - \hat{w}_i^t \geq w(1) - \alpha_t - (w(i) + \alpha_t) = w(1) - w(i) - 2\alpha_2,$$

we eliminate a suboptimal arm $i \neq 1$ when

$$w(1) - w(i) - 2\alpha_t = \Delta_i - 2\alpha_t \geq 2\alpha_t.$$

Step 1. Concentration inequality:

$$\Pr\left(\cup_{i\in[L]}\cup_{t\in\mathbb{N}}\left\{|\hat{w}_i^t - w(i)| > \alpha_t\right\}\right) \leq \delta.$$

Step 2. Assume $\cap_{i\in[L]}\cap_{t\in\mathbb{N}}\left\{|\hat{w}_i^t - w(i)| \leq \alpha_t\right\}$ holds. Recall that we eliminate arm i when $\hat{w}_{\max}^t - \hat{w}_i^t \geq 2\alpha_t$. Since

$$\hat{w}_{\max}^t - \hat{w}_i^t \geq \hat{w}_1^t - \hat{w}_i^t \geq w(1) - \alpha_t - (w(i) + \alpha_t) = w(1) - w(i) - 2\alpha_2,$$

we eliminate a suboptimal arm $i \neq 1$ when

$$w(1) - w(i) - 2\alpha_t = \Delta_i - 2\alpha_t \geq 2\alpha_t.$$

Step 3. When each arm has been sampled for

$$t_i = O\left(\frac{\log(L/(\delta\Delta_i))}{\Delta_i^2}\right)$$

times, we have $\alpha_t \leq \Delta_i/4$ and arm i will be eliminated.

Hence, the time complexity would be

$$t_2 + \sum_{i=2}^{L} t_i = O\left(\sum_{t=1}^{L}\frac{\log(L/(\delta\Delta_i))}{\Delta_i^2}\right) = \tilde{O}(H_1), \quad H_1 = \sum_{i=1}^{L}\frac{1}{\Delta_i^2} \text{ (hardness)}.$$

Median Elimination(ϵ, δ) (Even-Dar et al., 2002)

♠ With probability $1-\delta$, identify an ϵ-optimal arm i: $w(i) \geq \max_{j \in [L]} w(j) - \epsilon$.

Algorithm 2: Median Elimination(ϵ, δ) (Even-Dar et al., 2002)

♠ With probability $1-\delta$, identify an ϵ-optimal arm i: $w(i) \geq \max_{j \in [L]} w(j) - \epsilon$.

Algorithm 2: MEDIAN ELIMINATION(ϵ, δ) (Even-Dar et al., 2002)

1: Input: Survival set $S = [L]$. Set $\epsilon_1 = \epsilon/4$, $\delta_1 = \delta/2$, $\ell = 1$.

2: Sample each arm $i \in S$ for $\frac{1}{(\epsilon_\ell/2)^2} \log(3/\delta_\ell)$ times, and let $\hat{w}_i^t$ denote its average reward.

♠ With probability $1-\delta$, identify an ϵ-optimal arm i: $w(i) \geq \max_{j \in [L]} w(j) - \epsilon$.

Algorithm 2: MEDIAN ELIMINATION(ϵ, δ) (Even-Dar et al., 2002)

1: Input: Survival set $S = [L]$. Set $\epsilon_1 = \epsilon/4$, $\delta_1 = \delta/2$, $\ell = 1$.
2: Sample each arm $i \in S$ for $\frac{1}{(\epsilon_\ell/2)^2} \log(3/\delta_\ell)$ times, and let $\hat{w}_i^t$ denote its average reward.
3: Find the median of $\hat{w}_i^\ell$, denoted by $m_\ell := \text{median}\big(\{\hat{w}_i^\ell : i \in S_\ell\}\big)$.
4: Let $S_{\ell+1} = S_\ell \setminus \{i : \hat{w}_i^\ell < m_\ell\}$.

♠ With probability $1-\delta$, identify an ϵ-optimal arm i: $w(i) \geq \max_{j \in [L]} w(j) - \epsilon$.

Algorithm 2: MEDIAN ELIMINATION(ϵ, δ) (Even-Dar et al., 2002)

1: Input: Survival set $S = [L]$. Set $\epsilon_1 = \epsilon/4$, $\delta_1 = \delta/2$, $\ell = 1$.
2: Sample each arm $i \in S$ for $\frac{1}{(\epsilon_\ell/2)^2} \log(3/\delta_\ell)$ times, and let $\hat{w}_i^t$ denote its average reward.
3: Find the median of $\hat{w}_i^\ell$, denoted by $m_\ell := \text{median}\big(\{\hat{w}_i^\ell : i \in S_\ell\}\big)$.
4: Let $S_{\ell+1} = S_\ell \setminus \{i : \hat{w}_i^\ell < m_\ell\}$.
5: $t = t + 1$.
6: If $|S| = 1$, Then output S.
Else $\epsilon_{\ell+1} = \frac{3}{4}\epsilon_\ell$, $\delta_{\ell+1} = \delta_\ell/2$, $\ell = \ell + 1$; Go to Step 2.

♠ With probability $1-\delta$, identify an ϵ-optimal arm i: $w(i) \geq \max_{j \in [L]} w(j) - \epsilon$.

Algorithm 2: MEDIAN ELIMINATION(ϵ, δ) (Even-Dar et al., 2002)

1: Input: Survival set $S = [L]$. Set $\epsilon_1 = \epsilon/4$, $\delta_1 = \delta/2$, $\ell = 1$.
2: Sample each arm $i \in S$ for $\frac{1}{(\epsilon_\ell/2)^2} \log(3/\delta_\ell)$ times, and let $\hat{w}_i^t$ denote its average reward.
3: Find the median of $\hat{w}_i^\ell$, denoted by $m_\ell := \text{median}\big(\{\hat{w}_i^\ell : i \in S_\ell\}\big)$.
4: Let $S_{\ell+1} = S_\ell \setminus \{i : \hat{w}_i^\ell < m_\ell\}$.
5: $t = t + 1$.
6: If $|S| = 1$, Then output S.
Else $\epsilon_{\ell+1} = \frac{3}{4}\epsilon_\ell$, $\delta_{\ell+1} = \delta_\ell/2$, $\ell = \ell + 1$; Go to Step 2.

Applying the same concentration inequality, we can show the time complexity of MEDIAN ELIMINATION(ϵ, δ) is

$$O\left(\frac{L \log(1/\delta)}{\epsilon^2}\right).$$

Lower bound (Garivier and Kaufmann, 2016)

For any δ-PAC algorithm and any bandit instance μ,

$$\mathbb{E}_\mu[\tau_\delta] \geq T^*(\mu) \log\left(\frac{4}{\delta}\right)$$

where

$$T^*(\mu)^{-1} := \sup_{w \in \Sigma_L} \inf_{\lambda \in \mathrm{Alt}(\mu)} \left(\sum_{i=1}^{L} w_i d(\mu_i, \lambda_i) \right).$$

Lower bound (Garivier and Kaufmann, 2016)

For any δ-PAC algorithm and any bandit instance μ,

$$\mathbb{E}_\mu[\tau_\delta] \geq T^*(\mu) \log\left(\frac{4}{\delta}\right)$$

where

$$T^*(\mu)^{-1} := \sup_{w \in \Sigma_L} \inf_{\lambda \in \mathrm{Alt}(\mu)} \left(\sum_{i=1}^{L} w_i d(\mu_i, \lambda_i) \right).$$

- For any instance $\mu = (\mu_1, \ldots, \mu_L) \in \mathcal{S}$
  - $\mathcal{S} = \{(\mu_1, \ldots, \mu_L) : \exists i^*(\mu) \in [L] \; s.t. \; \mu_{i^*(\mu)} > \mu_i \quad \forall i \neq i^*(\mu)\}$
  - Unique optimal arm: $i^*(\mu) = \arg\max_{i \in [L]} \mu_i$
  - "Alternative set": $\mathrm{Alt}(\mu) := \{\lambda \in \mathcal{S} : i^*(\lambda) \neq i^*(\mu)\}$
- Set of probability distributions on $[L]$

$$\Sigma_L = \left\{ (w_1, \ldots, w_L) \in (0,1]^L : \sum_{i=1}^{L} w_i = 1 \right\}$$

Proof strategy of lower bound

- Let $\lambda \in \mathrm{Alt}(\mu)$ and define event $E = \{\tau_\delta < \infty, i_{\mathrm{out}}(\mu) \neq i^*(\lambda)\} \in \mathcal{F}_{\tau_\delta}$. Then

$$
\begin{aligned}
2\delta &\geq \mathbb{P}_\mu(\tau_\delta < \infty \text{ and } i_{\mathrm{out}}(\mu) \neq i^*(\mu)) + \mathbb{P}_\mu(\tau_\delta < \infty \text{ and } i_{\mathrm{out}}(\mu) \neq i^*(\lambda)) \\
&\geq \mathbb{P}_\mu(E^c) + \mathbb{P}_\lambda(E) \\
&\geq \frac{1}{2} \exp\left(- \sum_{i=1}^{L} \mathbb{E}_\mu[T_i(\tau_\delta)] D(\mu_i, \lambda_i) \right). \qquad \text{Bretagnolle–Huber inequality}
\end{aligned}
$$

Proof strategy of lower bound

- Let $\lambda \in \mathrm{Alt}(\mu)$ and define event $E = \{\tau_\delta < \infty, i_{\mathrm{out}}(\mu) \neq i^*(\lambda)\} \in \mathcal{F}_{\tau_\delta}$. Then

$$\begin{aligned} 2\delta &\geq \mathbb{P}_\mu(\tau_\delta < \infty \text{ and } i_{\mathrm{out}}(\mu) \neq i^*(\mu)) + \mathbb{P}_\mu(\tau_\delta < \infty \text{ and } i_{\mathrm{out}}(\mu) \neq i^*(\lambda)) \\ &\geq \mathbb{P}_\mu(E^c) + \mathbb{P}_\lambda(E) \\ &\geq \frac{1}{2} \exp\left(- \sum_{i=1}^{L} \mathbb{E}_\mu[T_i(\tau_\delta)] D(\mu_i, \lambda_i) \right). \end{aligned}$$

Bretagnolle–Huber inequality

- Rearranging,

$$\sum_{i=1}^{L} \mathbb{E}_\mu[T_i(\tau_\delta)] D(\mu_i, \lambda_i) \geq \log \frac{4}{\delta}$$

Proof strategy of lower bound

- Let $\lambda \in \mathrm{Alt}(\mu)$ and define event $E = \{\tau_\delta < \infty, i_{\mathrm{out}}(\mu) \neq i^*(\lambda)\} \in \mathcal{F}_{\tau_\delta}$. Then

$$\begin{aligned} 2\delta &\geq \mathbb{P}_\mu(\tau_\delta < \infty \text{ and } i_{\mathrm{out}}(\mu) \neq i^*(\mu)) + \mathbb{P}_\mu(\tau_\delta < \infty \text{ and } i_{\mathrm{out}}(\mu) \neq i^*(\lambda)) \\ &\geq \mathbb{P}_\mu(E^c) + \mathbb{P}_\lambda(E) \\ &\geq \frac{1}{2}\exp\left(-\sum_{i=1}^{L} \mathbb{E}_\mu[T_i(\tau_\delta)] D(\mu_i, \lambda_i)\right). \end{aligned}$$

Bretagnolle–Huber inequality

- Rearranging,

$$\sum_{i=1}^{L} \mathbb{E}_\mu[T_i(\tau_\delta)] D(\mu_i, \lambda_i) \geq \log \frac{4}{\delta}$$

- Using this and the definition of $T^*(\mu)$,

$$\begin{aligned} \frac{\mathbb{E}_\mu[\tau_\delta]}{T^*(\mu)} &= \mathbb{E}_\mu[\tau_\delta] \sup_{w \in \Sigma_L} \inf_{\lambda \in \mathrm{Alt}(\mu)} \sum_{i=1}^{L} w_i D(\mu_i, \lambda_i) \\ &\geq \cancel{\mathbb{E}_\mu[\tau_\delta]} \inf_{\lambda \in \mathrm{Alt}(\mu)} \sum_{i=1}^{L} \frac{\mathbb{E}_\mu[T_i(\tau_\delta)]}{\cancel{\mathbb{E}_\mu[\tau_\delta]}} D(\mu_i, \lambda_i) \\ &= \inf_{\lambda \in \mathrm{Alt}(\mu)} \sum_{i=1}^{L} \mathbb{E}_\mu[T_i(\tau_\delta)] D(\mu_i, \lambda_i) \geq \log \frac{4}{\delta}. \end{aligned}$$

Proof strategy of lower bound

- Let $\lambda \in \mathrm{Alt}(\mu)$ and define event $E = \{\tau_\delta < \infty, i_{\mathrm{out}}(\mu) \neq i^*(\lambda)\} \in \mathcal{F}_{\tau_\delta}$. Then

$$\begin{aligned} 2\delta &\geq \mathbb{P}_\mu(\tau_\delta < \infty \text{ and } i_{\mathrm{out}}(\mu) \neq i^*(\mu)) + \mathbb{P}_\mu(\tau_\delta < \infty \text{ and } i_{\mathrm{out}}(\mu) \neq i^*(\lambda)) \\ &\geq \mathbb{P}_\mu(E^c) + \mathbb{P}_\lambda(E) \\ &\geq \frac{1}{2}\exp\left(-\sum_{i=1}^{L} \mathbb{E}_\mu[T_i(\tau_\delta)] D(\mu_i, \lambda_i)\right). \end{aligned}$$

Bretagnolle–Huber inequality

- Rearranging,

$$\sum_{i=1}^{L} \mathbb{E}_\mu[T_i(\tau_\delta)] D(\mu_i, \lambda_i) \geq \log \frac{4}{\delta}$$

- Using this and the definition of $T^*(\mu)$,

$$\begin{aligned} \frac{\mathbb{E}_\mu[\tau_\delta]}{T^*(\mu)} &= \mathbb{E}_\mu[\tau_\delta] \sup_{w \in \Sigma_L} \inf_{\lambda \in \mathrm{Alt}(\mu)} \sum_{i=1}^{L} w_i D(\mu_i, \lambda_i) \\ &\geq \cancel{\mathbb{E}_\mu[\tau_\delta]} \inf_{\lambda \in \mathrm{Alt}(\mu)} \sum_{i=1}^{L} \frac{\mathbb{E}_\mu[T_i(\tau_\delta)]}{\cancel{\mathbb{E}_\mu[\tau_\delta]}} D(\mu_i, \lambda_i) \\ &= \inf_{\lambda \in \mathrm{Alt}(\mu)} \sum_{i=1}^{L} \mathbb{E}_\mu[T_i(\tau_\delta)] D(\mu_i, \lambda_i) \geq \log \frac{4}{\delta}. \end{aligned}$$

We thus have the asymptotic lower bound on the time complexity:

$$\liminf_{\delta \to 0} \frac{\mathbb{E}_\mu[\tau_\delta]}{\log(1/\delta)} \geq T^*(\mu).$$

We thus have the asymptotic lower bound on the time complexity:

$$\liminf_{\delta \to 0} \frac{\mathbb{E}_\mu[\tau_\delta]}{\log(1/\delta)} \geq T^*(\mu).$$

A matching upper bound can be achieved by TRACK & STOP

$$\mathbb{P}_\mu \left(\limsup_{\delta \to 0} \frac{\tau_\delta}{\log(1/\delta)} \leq T^*(\mu) \right) = 1,$$

or

$$\limsup_{\delta \to 0} \frac{\mathbb{E}_\mu[\tau_\delta]}{\log(1/\delta)} \leq T^*(\mu).$$

Algorithm 3: TRACK & STOP (Garivier and Kaufmann, 2016)

1: Let $N_i(t) = \sum_{u=1}^{t} 1\{S_u = i\}$ be the number of pulls of arm i,

$$\hat{\mu}_i(t) = \frac{1}{N_i(t)} \sum_{u=1}^{t} W_t(i) 1\{S_u = i\}$$ be the empirical mean of arm i.

Set $\hat{\mu}(t) = (\hat{\mu}_1(t), \hat{\mu}_2(t), \ldots, \hat{\mu}_L(t))$.

2: Sample each arm once and update $t = L$, $N_i(L)$, $\hat{\mu}_i(L)$.

Algorithm 3: TRACK & STOP (Garivier and Kaufmann, 2016)

Let $N_i(t) = \sum_{u=1}^{t} 1\{S_u = i\}$ be the number of pulls of arm i,

$\hat{\mu}_i(t) = \frac{1}{N_i(t)} \sum_{u=1}^{t} W_t(i) 1\{S_u = i\}$ be the empirical mean of arm i.

Set $\hat{\mu}(t) = (\hat{\mu}_1(t), \hat{\mu}_2(t), \ldots, \hat{\mu}_L(t))$.

Sample each arm once and update $t = L$, $N_i(L)$, $\hat{\mu}_i(L)$.

while *Stopping condition (Generalized Likelihood Ratio statistic)* is not satisfied **do**

  Sample arm S_{t+1} by C-Tracking/D-Tracking rule.

  Let $t = t + 1$, and update $N_i(t)$, $\hat{\mu}_i(t)$.

end while

Output $\hat{i} = \arg\max_{i \in [L]} \hat{\mu}_i(t)$.

Sampling rule

C-Tracking: $S_{t+1} \in \arg\max_{i \in [L]} \sum_{\tau=0}^{t} w_i^{\epsilon_\tau}(\hat{\mu}(\tau)) - N_i(t)$

D-Tracking: $S_{t+1} \in \begin{cases} \arg\min_{i \in U_t} N_i(t) & \text{if } U_t \neq \emptyset \quad \text{(forced exploration)} \\ \arg\max_{i \in [L]} t w_i^{\epsilon_t}(\hat{\mu}(t)) - N_i(t) & \text{else} \quad \text{(directed tracking)} \end{cases}$

Sampling rule

C-Tracking: $S_{t+1} \in \arg\max_{i\in[L]} \sum_{\tau=0}^{t} w_i^{\epsilon_\tau}(\hat{\mu}(\tau)) - N_i(t)$

D-Tracking: $S_{t+1} \in \begin{cases} \arg\min_{i\in U_t} N_i(t) & \text{if } U_t \neq \emptyset \quad \text{(forced exploration)} \\ \arg\max_{i\in[L]} t w_i^{\epsilon_t}(\hat{\mu}(t)) - N_i(t) & \text{else} \quad \text{(directed tracking)} \end{cases}$

$$w^*(\mu) = \arg\max_{w\in\Sigma_L} \inf_{\lambda\in \mathsf{Alt}(\mu)} \Big(\sum_{i=1}^{L} w_i d(w_i, \lambda_i) \Big),$$

- Proportion of arm draws of any strategy matches the lower bound

$\epsilon_t = (L^2 + t)^{-1/2}/2,$

$w^\epsilon(\mu)$: L^∞ projection of $w^*(\mu)$ onto $\Sigma_L^{(\epsilon)} = \Big\{(w_1, \ldots, w_L) \in [\epsilon, 1]^L : \sum_{i=1}^{L} = 1\Big\}$

Sampling rule

C-Tracking: $S_{t+1} \in \arg\max_{i\in[L]} \sum_{\tau=0}^{t} w_i^{\epsilon_\tau}(\hat{\mu}(\tau)) - N_i(t)$

D-Tracking: $S_{t+1} \in \begin{cases} \arg\min_{i\in U_t} N_i(t) & \text{if } U_t \neq \emptyset \quad \text{(forced exploration)} \\ \arg\max_{i\in[L]} t w_i^{\epsilon_t}(\hat{\mu}(t)) - N_i(t) & \text{else} \quad \text{(directed tracking)} \end{cases}$

$$w^*(\mu) = \arg\max_{w\in\Sigma_L} \inf_{\lambda\in \mathsf{Alt}(\mu)} \Big(\sum_{i=1}^{L} w_i d(w_i, \lambda_i) \Big),$$

- Proportion of arm draws of any strategy matches the lower bound

$\epsilon_t = (L^2 + t)^{-1/2}/2,$

$w^\epsilon(\mu)$: L^∞ projection of $w^*(\mu)$ onto $\Sigma_L^{(\epsilon)} = \Big\{(w_1, \ldots, w_L) \in [\epsilon, 1]^L : \sum_{i=1}^{L} = 1\Big\}$

Outline

1 What is multi-armed bandits (MAB)?

BAI: fixed-budget

Theoretical study

▲ Propose a BAI algorithm in a fixed time horizon and upper bound its failure probability

▼ Derive a lower bound on the failure probability of any algorithm

• Evaluate theoretical findings with experiments

BAI: fixed-budget

Theoretical study

▲ Propose a BAI algorithm in a fixed time horizon and upper bound its failure probability

▼ Derive a lower bound on the failure probability of any algorithm

- Evaluate theoretical findings with experiments

Simple pure exploration in stochastic bandits

- to identify the best arm with the largest mean: $i^* = \arg\max_{i \in [L]} w(i)$

BAI: fixed-budget

Theoretical study

▲ Propose a BAI algorithm in a fixed time horizon and upper bound its failure probability

▼ Derive a lower bound on the failure probability of any algorithm

• Evaluate theoretical findings with experiments

Simple pure exploration in stochastic bandits

• to identify the best arm with the largest mean: $i^* = \arg\max_{i \in [L]} w(i)$

♠ UCB-based

UCB-E(a) (Audibert and Bubeck, 2010)

♠ Successive elimination

Sequential Halving (Karnin et al., 2013)

Algorithm 4: UCB-E(a) (Audibert and Bubeck, 2010)

1: **Input:** time budget T, size of ground set of items L, parameter a.
2: For all $i \in [L]$, compute $N_{i,0}$, $\hat{w}_{i,0}$, $C_{i,0}$, $U_{i,0}$:

$$N_{i,t} = \sum_{u=1}^{t} 1\{i_u = i\}, \ \hat{w}_{i,t} = \frac{1}{N_{i,t}} \sum_{u=1}^{t} W_{i,t} \cdot 1\{i_u = i\},$$

$$C_{i,t} = \sqrt{\frac{a}{t}} \text{ if } t \geq 1, \qquad C_{i,0} = +\infty, \qquad U_{i,t} = \hat{g}_{i,t} + C_{i,t}.$$

3: **for** $t = 1, \ldots, T$ **do**
4: Pull item $i_t = \arg\max_{i \in [L]} U_{i,t-1}$.
5: Update $N_{i_t,t}$, $\hat{w}_{i_t,t}$, $C_{i,t}$, and $U_{i,t}$ for all i.
6: **end for**
7: Output $i_{\text{out}} = \arg\max_{i \in [L]} \hat{w}_{i,T}$.

Step 1: Concentration. Let $\mathcal{E}_i := \{\forall t \geq L, |\hat{w}_{i,t} - w(i)| \leq C_{i,t}/5\}$ for all $i \in [L]$. We apply concentration inequality to show that

$$\Pr\left(\bigcap_{i=1}^{L} \mathcal{E}_i\right) \geq 1 - 2TL\exp\left(-\frac{2a}{25}\right).$$

In the following, we prove that conditioned on the event $\bigcap_{i=1}^{L} \mathcal{E}_i$, we have $i_{\text{out}} = 1$, which concludes the proof.

We assume $\bigcap_{i=1}^{L} \mathcal{E}_i$ holds from now on. Since i_{out} is the item with the largest empirical mean, for all $i \neq i_{\text{out}}$, we have

$$\hat{w}_{i_{\text{out}},T} \geq \hat{w}_{i,t}, \quad \hat{w}_{i_{\text{out}},T} \geq w(i_{\text{out}}) - C_{i_{\text{out}},T}/5, \quad w(i) + C_{i,T}/5 \geq \hat{w}_{i,t}.$$

Consequently, to show $i_{\text{out}} = 1$, it is sufficient to show that

$$\frac{C_{i,T}}{5} \leq \frac{\Delta_i}{2} \;\Leftrightarrow\; N_{it} \geq \frac{4}{25}\frac{a}{\Delta_i^2} \quad \forall\, i \in [L]. \tag{1}$$

Step 1: Concentration. Let $\mathcal{E}_i := \{\forall t \geq L, |\hat{w}_{i,t} - w(i)| \leq C_{i,t}/5\}$ for all $i \in [L]$. We apply concentration inequality to show that

$$\Pr\left(\bigcap_{i=1}^{L} \mathcal{E}_i\right) \geq 1 - 2TL\exp\left(-\frac{2a}{25}\right).$$

In the following, we prove that conditioned on the event $\bigcap_{i=1}^{L} \mathcal{E}_i$, we have $i_{\text{out}} = 1$, which concludes the proof.

We assume $\bigcap_{i=1}^{L} \mathcal{E}_i$ holds from now on. Since i_{out} is the item with the largest empirical mean, for all $i \neq i_{\text{out}}$, we have

$$\hat{w}_{i_{\text{out}},T} \geq \hat{w}_{i,t}, \quad \hat{w}_{i_{\text{out}},T} \geq w(i_{\text{out}}) - C_{i_{\text{out}},T}/5, \quad w(i) + C_{i,T}/5 \geq \hat{w}_{i,t}.$$

Consequently, to show $i_{\text{out}} = 1$, it is sufficient to show that

$$\frac{C_{i,T}}{5} \leq \frac{\Delta_i}{2} \;\Leftrightarrow\; N_{it} \geq \frac{4}{25}\frac{a}{\Delta_i^2} \quad \forall\, i \in [L]. \tag{1}$$

Step 2: Upper bound $N_{i,T}$ $(i \neq 1)$. To begin with, we prove by induction that

$$N_{i,t} \leq \frac{36}{25}\frac{a}{\Delta_i^2} \quad \forall\, i \neq 1. \tag{2}$$

Step 3: Lower bound $N_{i,T}$ $(i \neq 1)$. Next, we again prove by induction that

$$N_{i,t} \geq \frac{4}{25} \min \left\{ \frac{a}{\Delta_i^2}, \frac{25}{36}(N_{1,t} - 1) \right\} \quad \forall i \neq 1. \tag{3}$$

Step 3: Lower bound $N_{i,T}$ $(i \neq 1)$. Next, we again prove by induction that

$$N_{i,t} \geq \frac{4}{25} \min \left\{ \frac{a}{\Delta_i^2}, \frac{25}{36}(N_{1,t} - 1) \right\} \quad \forall i \neq 1. \tag{3}$$

Step 4: Lower bound on $N_{1,T}$. Recall that we want to show (**??**). (i) To show (**??**) holds for all $i \neq 1$, (**??**) indicates that it is sufficient to show that

$$\frac{25}{36}(N_{1,t} - 1) \geq \frac{a}{\Delta_i^2} \quad \forall i \neq 1.$$

(ii) In order to show (**??**) holds for all $i = 1$, we apply (**??**), $t = \sum_{i=1}^{L} N_{i,t}$ and

$$\frac{36}{25} H_1 a \leq T - L \Leftrightarrow a \leq \frac{25(T-L)}{36 H_1}, \quad H_1 = \sum_{i=1}^{L} \frac{1}{\Delta_i^2}.$$

Step 3: Lower bound $N_{i,T}$ $(i \neq 1)$**.** Next, we again prove by induction that

$$N_{i,t} \geq \frac{4}{25} \min \left\{ \frac{a}{\Delta_i^2}, \frac{25}{36}(N_{1,t}-1) \right\} \quad \forall i \neq 1. \tag{3}$$

Step 4: Lower bound on $N_{1,T}$**.** Recall that we want to show (**??**). (i) To show (**??**) holds for all $i \neq 1$, (**??**) indicates that it is sufficient to show that

$$\frac{25}{36}(N_{1,t}-1) \geq \frac{a}{\Delta_i^2} \quad \forall i \neq 1.$$

(ii) In order to show (**??**) holds for all $i = 1$, we apply (**??**), $t = \sum_{i=1}^{L} N_{i,t}$ and

$$\frac{36}{25} H_1 a \leq T - L \;\Leftrightarrow a \leq \frac{25(T-L)}{36 H_1}, \quad H_1 = \sum_{i=1}^{L} \frac{1}{\Delta_i^2}.$$

Step 5: Conclusion. The failure probability is

$$2TL \exp\left(-\frac{2a}{25} \right) \quad \forall\, a \leq \frac{25(T-L)}{36 H_1}$$

and achieves the minimum,

$$2TL \exp\left(-\frac{T-L}{18 H_1} \right) \quad \text{when } a = \frac{25(T-L)}{36 H_1}.$$

Step 3: Lower bound $N_{i,T}$ $(i \neq 1)$**.** Next, we again prove by induction that

$$N_{i,t} \geq \frac{4}{25} \min \left\{ \frac{a}{\Delta_i^2}, \frac{25}{36}(N_{1,t}-1) \right\} \quad \forall i \neq 1. \tag{3}$$

Step 4: Lower bound on $N_{1,T}$**.** Recall that we want to show (**??**). (i) To show (**??**) holds for all $i \neq 1$, (**??**) indicates that it is sufficient to show that

$$\frac{25}{36}(N_{1,t}-1) \geq \frac{a}{\Delta_i^2} \quad \forall i \neq 1.$$

(ii) In order to show (**??**) holds for all $i = 1$, we apply (**??**), $t = \sum_{i=1}^{L} N_{i,t}$ and

$$\frac{36}{25} H_1 a \leq T - L \;\Leftrightarrow a \leq \frac{25(T-L)}{36 H_1}, \quad H_1 = \sum_{i=1}^{L} \frac{1}{\Delta_i^2}.$$

Step 5: Conclusion. The failure probability is

$$2TL \exp\left(-\frac{2a}{25} \right) \quad \forall\, a \leq \frac{25(T-L)}{36 H_1}$$

and achieves the minimum, however, requiring prior knowledge: hardness H_1

$$2TL \exp\left(-\frac{T-L}{18 H_1} \right) \quad \text{when } a = \frac{25(T-L)}{36 H_1}.$$

Algorithm 5: Sequantial Halving (SH) (Karnin et al., 2013)

1: Input: time budget T, size of ground set L.
2: Set $M = \lceil \log_2 L \rceil$, $N = \lfloor T/M \rfloor$, $T_0 = 0$, $A_0 = [L]$.

Algorithm 5: Sequantial Halving (SH) (Karnin et al., 2013)

1: Input: time budget T, size of ground set L.
2: Set $M = \lceil \log_2 L \rceil$, $N = \lfloor T/M \rfloor$, $T_0 = 0$, $A_0 = [L]$.

M : number of phases

N : length of each phase

T_m : last time step of phase m

A_m : active set after phase m

Algorithm 5: Sequantial Halving (SH) (Karnin et al., 2013)

Input: time budget T, size of ground set L.
Set $M = \lceil \log_2 L \rceil$, $N = \lfloor T/M \rfloor$, $T_0 = 0$, $A_0 = [L]$.
for phase $m = 1, 2, \ldots, M$ **do**
  Set $T_m = T_{m-1} + N, q_m = 1/|A_{m-1}|, n_m = \lfloor q_m N \rfloor$.
  for $t = T_{m-1} + 1, \ldots, T_m$ **do**
    Pull $i \in A_{m-1}$ with for n_m times in order and observe $W_t(i)$.
  end for

Algorithm 5: Sequantial Halving (SH) (Karnin et al., 2013)

Input: time budget T, size of ground set L.
Set $M = \lceil \log_2 L \rceil$, $N = \lfloor T/M \rfloor$, $T_0 = 0$, $A_0 = [L]$.
for phase $m = 1, 2, \ldots, M$ **do**
  Set $T_m = T_{m-1} + N, q_m = 1/|A_{m-1}|, n_m = \lfloor q_m N \rfloor$.
  for $t = T_{m-1} + 1, \ldots, T_m$ **do**
    Pull $i \in A_{m-1}$ with for n_m times in order and observe $W_t(i)$.
  end for
  For all $i \in A_{m-1}$, set

$$S_m(i) = \sum_{t=T_{m-1}+1}^{T_m} W_t(i_t) \cdot \mathbb{I}\{i_t = i\}, \ \hat{w}_m(i) = \frac{S_m(i)}{n_m}.$$

  Let A_m contain the $\lceil L/2^m \rceil$ items with the highest $\hat{w}_m(i)$'s in A_{m-1}.

Algorithm 5: SEQUANTIAL HALVING (SH) (Karnin et al., 2013)

Input: time budget T, size of ground set L.
Set $M = \lceil \log_2 L \rceil$, $N = \lfloor T/M \rfloor$, $T_0 = 0$, $A_0 = [L]$.
for phase $m = 1, 2, \ldots, M$ **do**
  Set $T_m = T_{m-1} + N, q_m = 1/|A_{m-1}|, n_m = \lfloor q_m N \rfloor$.
  for $t = T_{m-1} + 1, \ldots, T_m$ **do**
    Pull $i \in A_{m-1}$ with for n_m times in order and observe $W_t(i)$.
  end for
  For all $i \in A_{m-1}$, set

$$S_m(i) = \sum_{t=T_{m-1}+1}^{T_m} W_t(i_t) \cdot \mathbb{I}\{i_t = i\}, \ \hat{w}_m(i) = \frac{S_m(i)}{n_m}.$$

  Let A_m contain the $\lceil L/2^m \rceil$ items with the highest $\hat{w}_m(i)$'s in A_{m-1}.
end for
Output the single item $i_{\text{out}} \in A_M$.

Step 1: Assume that the best arm was not eliminated prior to phase m. Then

$$\Pr(\hat{w}_m(1) < \hat{w}_m(i)) \leq \exp\Big(-\frac{1}{2} n_m \Delta_i^2\Big) \quad \forall\, i \in S_m \setminus \{1\}.$$

Step 2: The probability that the best arm is eliminated in phase m is at most

$$3\exp\Big(-\frac{T}{8\log_2 L} \cdot \frac{\Delta_{i_m}^2}{i_m}\Big)$$

where $i_m = L/2^{m+2}$.

Step 3: The failure probability can be bounded as follows:

$$\begin{aligned} 3\sum_{m=1}^{\log_2 L} \exp\Big(-\frac{T}{8\log_2 L} \cdot \frac{\Delta_{i_m}^2}{i_m}\Big) &\leq 3\sum_{m=1}^{\log_2 L} \exp\Big(-\frac{T}{8\log_2 L} \cdot \frac{1}{\max_i i\Delta_i^{-2}}\Big) \\ &= O\Big(\log_2 L \exp\Big(-\frac{T}{8H_2\log_2 L}\Big)\Big) \end{aligned}$$

when the hardness is measured by

$$H_2 = \max_{i\in[L]} \frac{i}{\Delta_i^2}.$$

BAI: fixed-budget

Algorithm/Instance	Reference	Failure probability e_T
$\text{UCB-E}\left(\frac{25(T-L)}{36H_1}\right)$	Audibert and Bubeck (2010)	$2TL\exp\left(-\frac{T-L}{18H_1}\right)$
SR	Audibert and Bubeck (2010)	$L(L-1)\exp\left(-\frac{T-L}{(1/2+\sum_{i=2}^{L}1/i)H_2}\right)$
$\text{UGapEb}\left(\frac{T-L}{16H_2}\right)$	Gabillon et al. (2012)	$2TL\exp\left(-\frac{T-L}{8H_2}\right)$
SAR	Bubeck et al. (2013)	$2L^2\exp\left(-\frac{T-L}{8(1/2+\sum_{i=2}^{L}1/i)H_2}\right)$
SH	Karnin et al. (2013)	$3\log_2 L\cdot\exp\left(-\frac{T}{8H_1\log_2 L}\right)$
NSE(p)	Shahrampour et al. (2017)	$(L-1)\exp\left(-\frac{2(T-L)}{H'_pC_p}\right)$
Stochastic Bandits	Carpentier and Locatelli (2016)	$\frac{1}{6}\exp\left(-\frac{400T}{H_2\log L}\right)$ (Lower Bound)

Shahrampour et al. (2017): $H'_p := \max_{i\neq 1}\frac{i^p}{\Delta_i^2}, \quad C_p := 2^{-p}+\sum_{i=2}^{L} i^{-p}\ \forall p>0.$

BAI: fixed-budget

Algorithm/Instance	Reference	Failure probability e_T
$\mathrm{UCB\text{-}E}\left(\frac{25(T-L)}{36H_1}\right)$	Audibert and Bubeck (2010)	$2TL\exp\left(-\frac{T-L}{18H_1}\right)$
SH	Karnin et al. (2013)	$3\log_2 L\cdot\exp\left(-\frac{T}{8H_1\log_2 L}\right)$
$\mathrm{NSE}(p)$	Shahrampour et al. (2017)	$(L-1)\exp\left(-\frac{2(T-L)}{H_p' C_p}\right)$
Stochastic Bandits	Carpentier and Locatelli (2016)	$\frac{1}{6}\exp\left(-\frac{400T}{H_2\log L}\right)$ (Lower Bound)

$$H_2 := \sum_{i=1}\frac{1}{\Delta_{i^2}}, \quad H_1 := \max_{i\neq 1}\frac{i}{\Delta_i^2}, \quad H_p' := \max_{i\neq 1}\frac{i^p}{\Delta_i^2}, \quad C_p := 2^{-p} + \sum_{i=2}^{L} i^{-p}\ \forall p>0.$$

BAI: fixed-budget

Algorithm/Instance	Reference	Failure probability e_T
$\text{UCB-E}\left(\frac{25(T-L)}{36H_1}\right)$	Audibert and Bubeck (2010)	$2TL\exp\left(-\frac{T-L}{18H_1}\right)$
SH	Karnin et al. (2013)	$3\log_2 L \cdot \exp\left(-\frac{T}{8H_1\log_2 L}\right)$
$\text{NSE}(p)$	Shahrampour et al. (2017)	$(L-1)\exp\left(-\frac{2(T-L)}{H_p' C_p}\right)$
Stochastic Bandits	Carpentier and Locatelli (2016)	$\frac{1}{6}\exp\left(-\frac{400T}{H_2\log L}\right)$ (Lower Bound)

$$H_2 := \sum_{i=1} \frac{1}{\Delta_{i^2}}, \quad H_1 := \max_{i\neq 1} \frac{i}{\Delta_i^2}, \quad H_p' := \max_{i\neq 1} \frac{i^p}{\Delta_i^2}, \quad C_p := 2^{-p} + \sum_{i=2}^{L} i^{-p} \ \forall p > 0.$$

- $H_2 \leq H_1 \leq H_2 \log(2L)$ (Audibert and Bubeck, 2010)
- Whether SH or $\text{NSE}(p)$ performs better depends on the instance, and SH does not involve a tunable parameter

Stochastic bandits

- Each arm $i \in [L]$ is associated with an unknown distribution $\nu(i)$, mean $w(i)$, and variance $\sigma(i)^2$.
- $\{W_t(i)\}_{t=1}^T$ is the i.i.d. sequence of rewards associated with arm i during the T time steps.

♠ **Question from real life:** do we always have i.i.d. data in real life?

BAI: fixed-budget + corruptions (Zhong et al., 2021)

Stochastic bandits

- Each arm $i \in [L]$ is associated with an unknown distribution $\nu(i)$, mean $w(i)$, and variance $\sigma(i)^2$.
- $\{W_t(i)\}_{t=1}^T$ is the i.i.d. sequence of rewards associated with arm i during the T time steps.

♠ **Question from real life:** do we always have i.i.d. data in real life?

$\Rightarrow$ Stochastic bandits with adversarial corruptions

▲ Propose algorithms with near-optimal performance guarantees

▼ Demonstrate (near-)optimality by designing an appropriate corruption strategy

Case 1: Biases and Contaminations in Clinical Trials

- To test the efficacy of a medicine on randomly chosen patients.

Case 1: Biases and Contaminations in Clinical Trials

- To test the efficacy of a medicine on randomly chosen patients.
- Possible biases and errors:

Case 1: Biases and Contaminations in Clinical Trials

- To test the efficacy of a medicine on randomly chosen patients.
- Possible biases and errors:
  - Loss-to-follow-up,

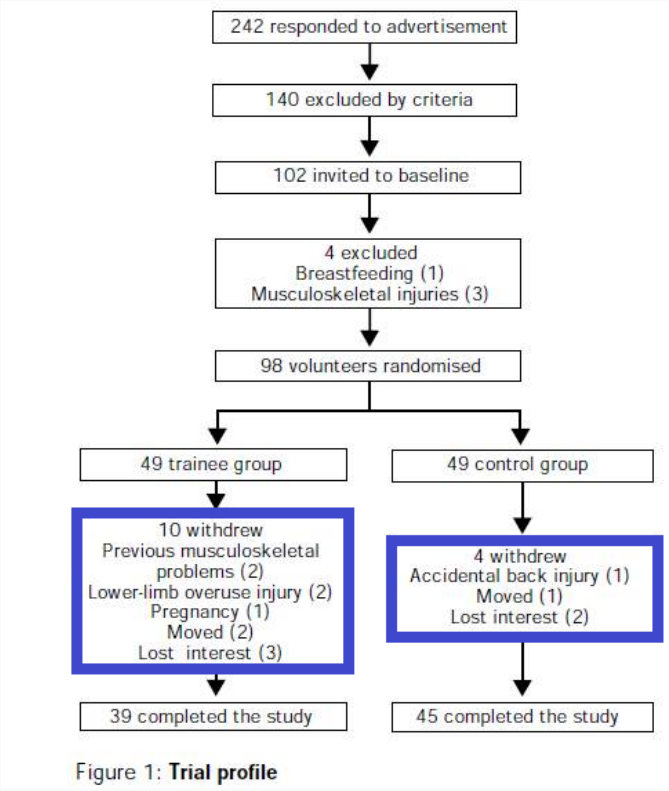

Figure 1: Loss-to-follow-up, boxed in blue.

Case 1: Biases and Contaminations in Clinical Trials

- To test the efficacy of a medicine on randomly chosen patients.
- Possible biases and errors:
  - Loss-to-follow-up,
  - Non-compliance. . .

Case 1: Biases and Contaminations in Clinical Trials

- To test the efficacy of a medicine on randomly chosen patients.
- Possible biases and errors:
  - Loss-to-follow-up,
  - Non-compliance. . .
- A lesson from COVID:
  - Not enough time to ensure i.i.d. samples!

Case 1: Biases and Contaminations in Clinical Trials

- To test the efficacy of a medicine on randomly chosen patients.
- Possible biases and errors:
  - Loss-to-follow-up,
  - Non-compliance. . .
- A lesson from COVID:
  - Not enough time to ensure i.i.d. samples!
- **How to identify the best medicine with *contaminated* data?**

Case 2: Fake Users in Online Recommendation Systems

- Paid reviews:
  - A major problem for recommender systems.

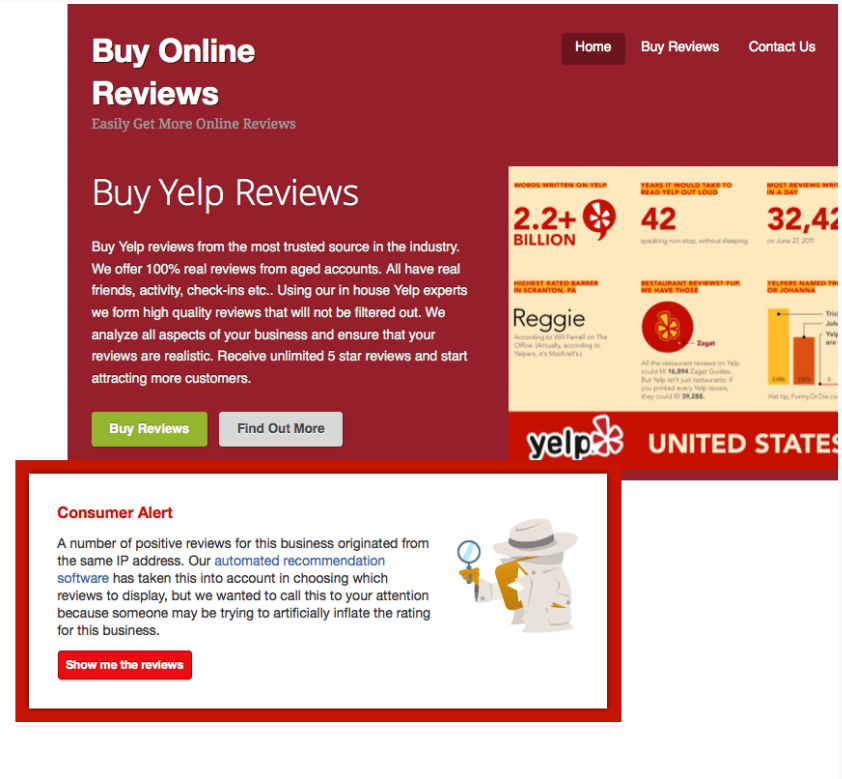

Figure 2: Buying fake reviews, and warnings about fake reviews.

Case 2: Fake Users in Online Recommendation Systems

- Paid reviews:
  - A major problem for recommender systems.
- Much effort to remove fake reviews.

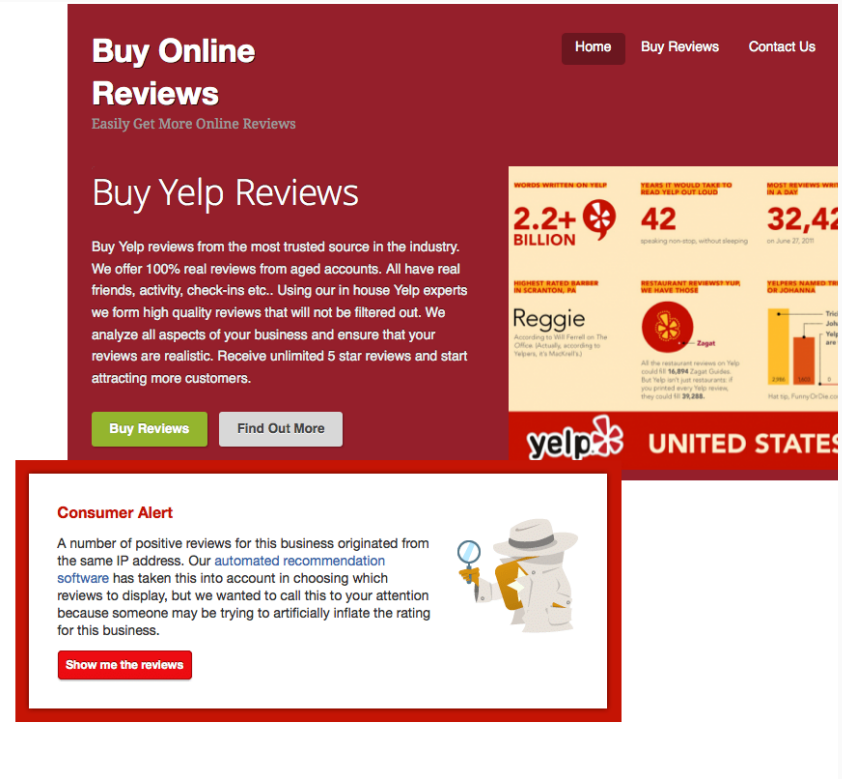

Figure 2: Buying fake reviews, and warnings about fake reviews.

Case 2: Fake Users in Online Recommendation Systems

- Paid reviews:
  - A major problem for recommender systems.
- Much effort to remove fake reviews.
- No fool-proof solution, anyone can review.

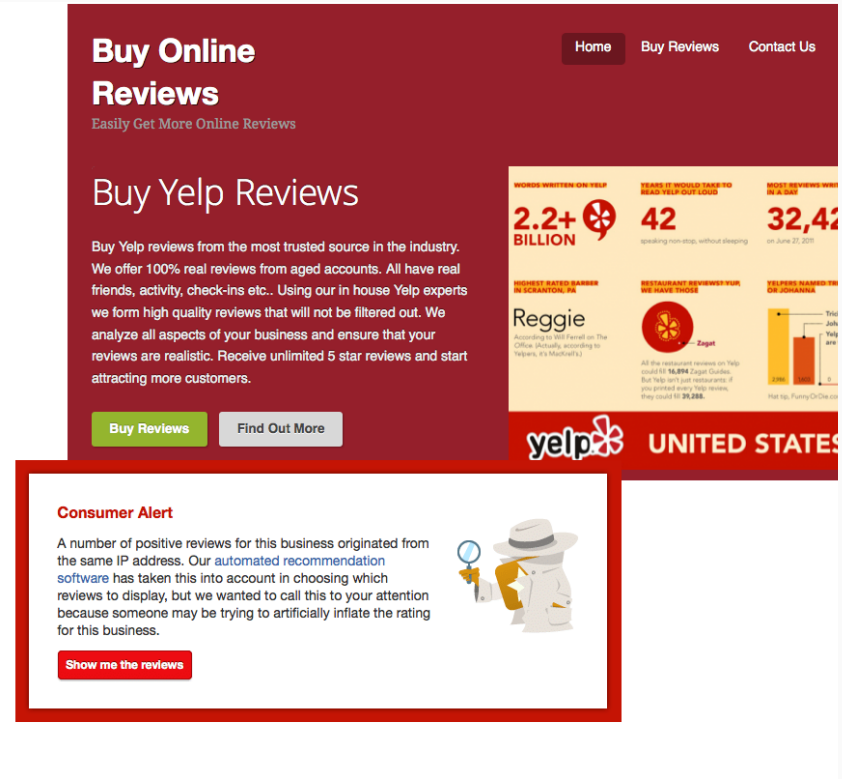

Figure 2: Buying fake reviews, and warnings about fake reviews.

Case 2: Fake Users in Online Recommendation Systems

- Paid reviews:
  - A major problem for recommender systems.
- Much effort to remove fake reviews.
- No fool-proof solution, anyone can review.
- **How to identify the best restaurants with contaminated data?**

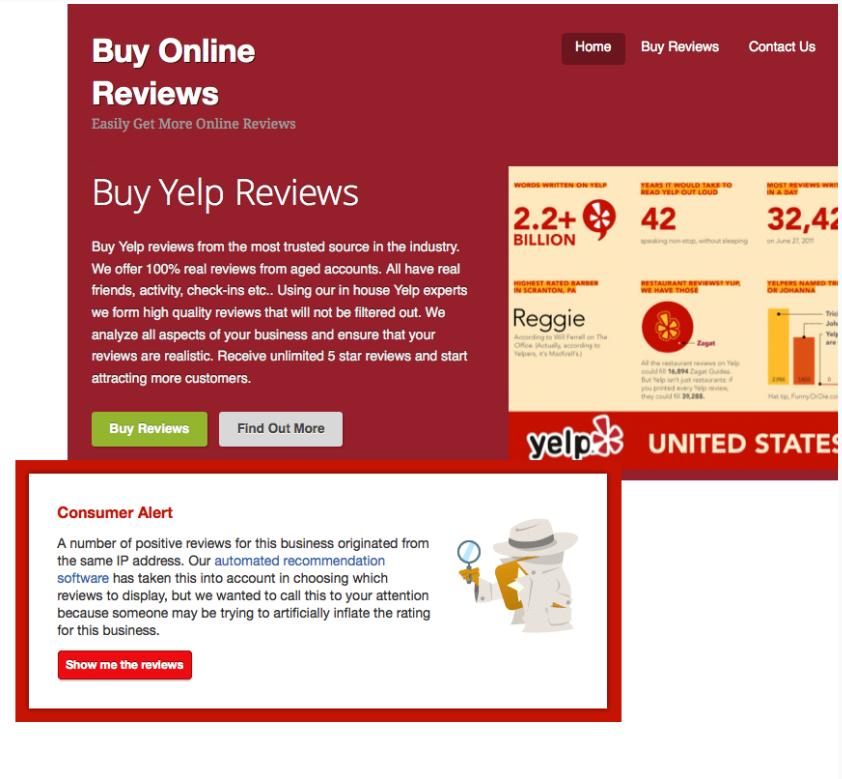

Figure 2: Buying fake reviews, and warnings about fake reviews.

Dynamic (Gupta et al., 2019)

- Ground set of L items indexed by $[L] := \{1, \dots, L\}$.
- Each item $i \in [L]$ is associated with an unknown mean $w(i) \in (0, 1]$.

Dynamic (Gupta et al., 2019)

- Ground set of L items indexed by $[L] := \{1, \ldots, L\}$.
- Each item $i \in [L]$ is associated with an unknown mean $w(i) \in (0, 1]$.
- Amount of adversarial corruptions is bounded by the unknown corruption budget C:

$$\sum_{t=1}^{T} \max_{i \in [L]} |c_t(i)| \leq C.$$

Dynamic (Gupta et al., 2019)

- Ground set of L items indexed by $[L] := \{1, \dots, L\}$.
- Each item $i \in [L]$ is associated with an unknown mean $w(i) \in (0, 1]$.
- Amount of adversarial corruptions is bounded by the unknown corruption budget C:

$$\sum_{t=1}^{T} \max_{i \in [L]} |c_t(i)| \leq C.$$

♠ At each time step $t = 1, \dots, T$:

1. A stochastic reward $W_t(i) \in [0, 1]$ is i.i.d. drawn for each item i.
2. The adversary observes $\{W_t(i)\}_{i \in [L]}$, and corrupts each $W_t(i)$ with $c_t(i) \in [-1, 1]$ if the corruption budget has not been depleted:

$$\tilde{W}_t(i) = W_t(i) + c_t(i) \in [0, 1]$$

3. The agent pulls $i_t \in [L]$ and observes the corrupted reward $\tilde{W}_t(i_t)$.

Dynamic (Gupta et al., 2019)

- Ground set of L items indexed by $[L] := \{1, \ldots, L\}$.
- Each item $i \in [L]$ is associated with an unknown mean $w(i) \in (0, 1]$.
- Amount of adversarial corruptions is bounded by the unknown corruption budget C:

$$\sum_{t=1}^{T} \max_{i \in [L]} |c_t(i)| \leq C.$$

♠ At each time step $t = 1, \ldots, T$:

1. A stochastic reward $W_t(i) \in [0, 1]$ is i.i.d. drawn for each item i.
2. The adversary observes $\{W_t(i)\}_{i \in [L]}$, and corrupts each $W_t(i)$ with $c_t(i) \in [-1, 1]$ if the corruption budget has not been depleted:

$$\tilde{W}_t(i) = W_t(i) + c_t(i) \in [0, 1]$$

3. The agent pulls $i_t \in [L]$ and observes the corrupted reward $\tilde{W}_t(i_t)$.

♠ At the end, the agent returns $i_{\text{out}} \in [L]$ as the recommendation.

Objective

- Assume $w(1) > w(2) \geq \ldots \geq w(L)$.
- Optimality gap of item i is $\Delta_{1,i} := w(1) - w(i)$.

Objective

- Assume $w(1) > w(2) \geq \ldots \geq w(L)$.
- Optimality gap of item i is $\Delta_{1,i} := w(1) - w(i)$.
- For fixed $\epsilon_C, \delta \in (0,1)$, an algorithm is said to be (ϵ_C, δ)-PAC (probably approximately correct) if

$$\mathbb{P}\left[\Delta_{1,i^{\pi,T}_{\text{out}}} > \epsilon_C\right] \leq \delta.$$

Objective

- Assume $w(1) > w(2) \geq \ldots \geq w(L)$.
- Optimality gap of item i is $\Delta_{1,i} := w(1) - w(i)$.
- For fixed $\epsilon_C, \delta \in (0,1)$, an algorithm is said to be (ϵ_C, δ)-PAC (probably approximately correct) if

$$\mathbb{P}\left[\Delta_{1,i^{\pi,T}_{\text{out}}} > \epsilon_C\right] \leq \delta.$$

♠ **Goal**: design an (ϵ_C, δ)-PAC algorithm π with both ϵ_C and δ small.

– $\epsilon_C < \Delta_{1,2}$: an (ϵ_C, δ)-PAC algorithm identifies the optimal item with probability at least $1 - \delta$.

Probabilistic Sequential Shrinking **Algorithm**

T time steps

Probabilistic Sequential Shrinking **Algorithm**

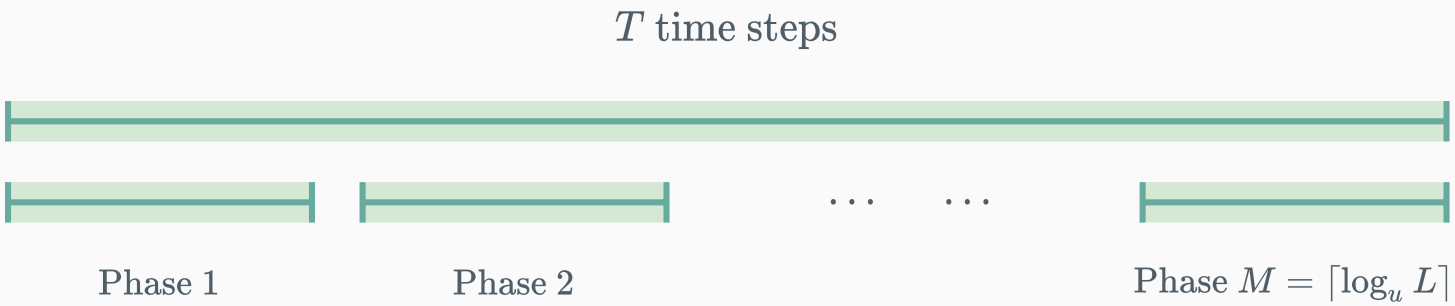

Probabilistic Sequential Shrinking **Algorithm**

Active set A_0

$A_0 = [L]$

Probabilistic Sequential Shrinking **Algorithm**

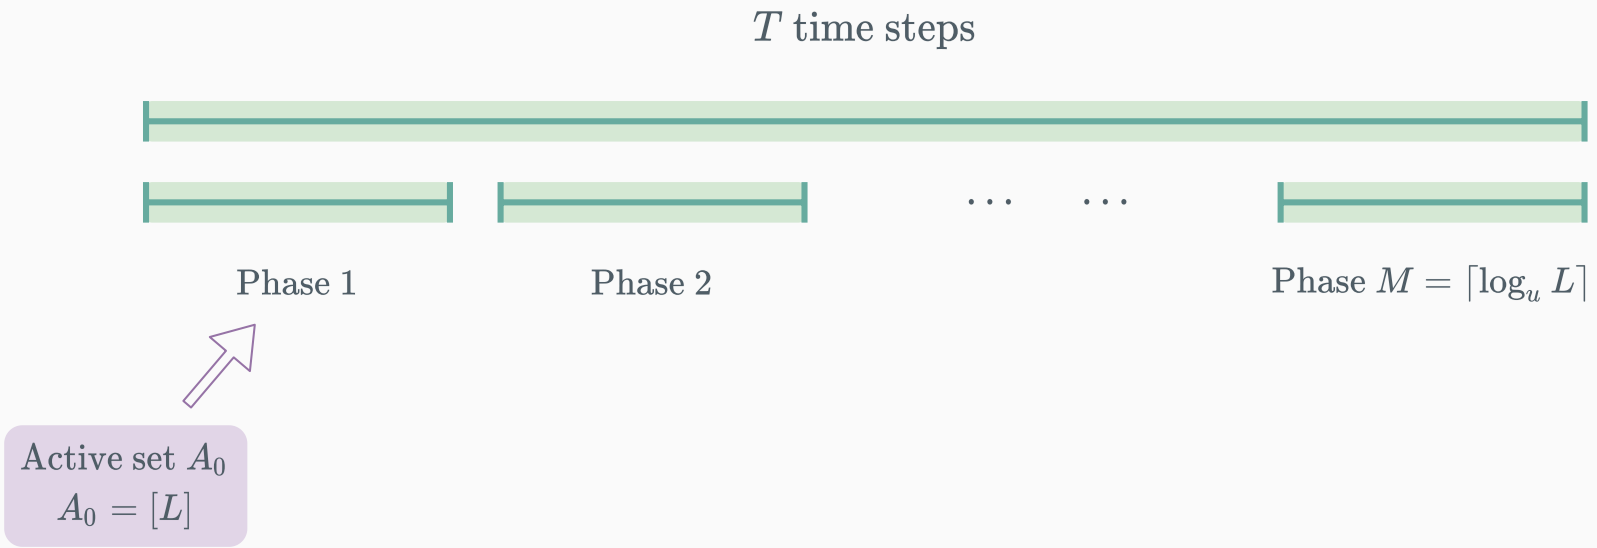

Probabilistic Sequential Shrinking **Algorithm**

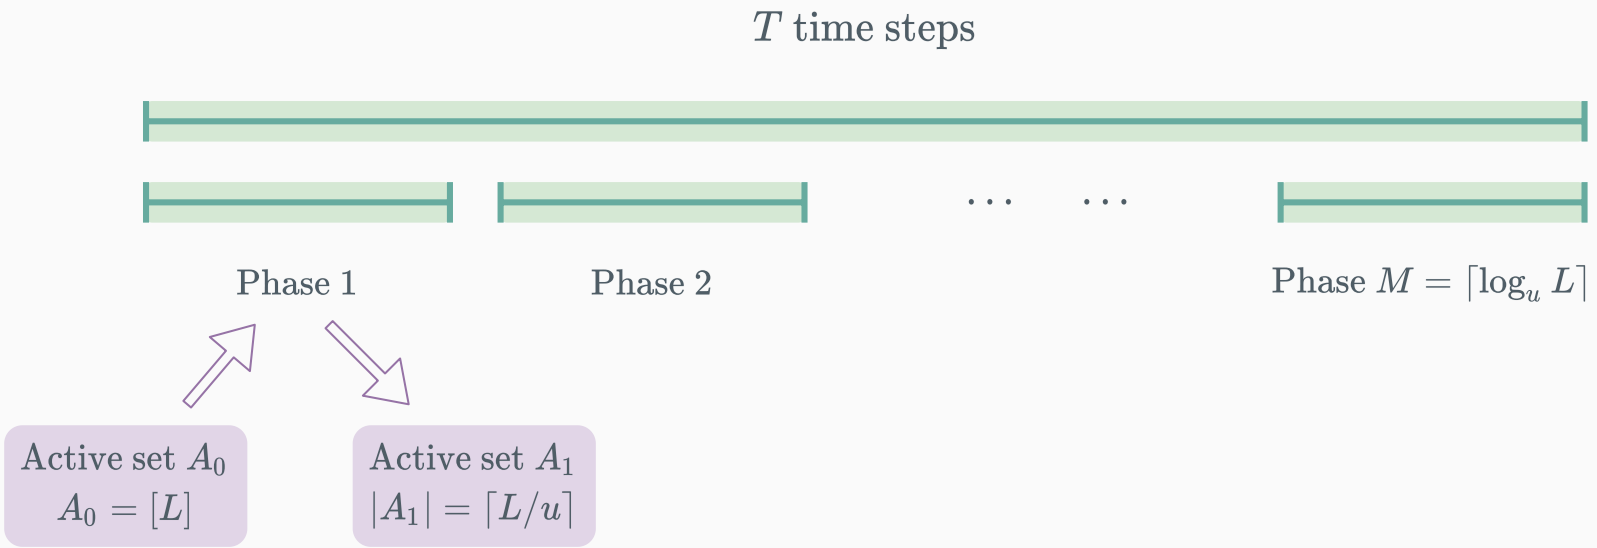

Probabilistic Sequential Shrinking **Algorithm**

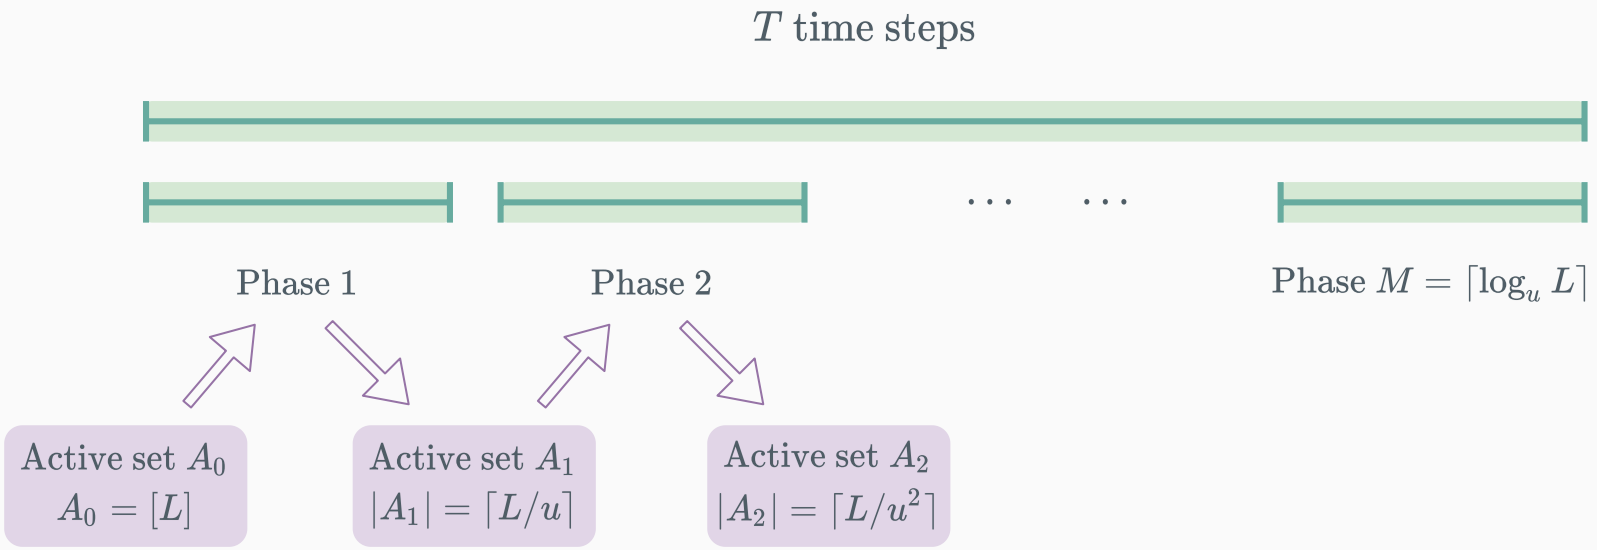

Probabilistic Sequential Shrinking **Algorithm**

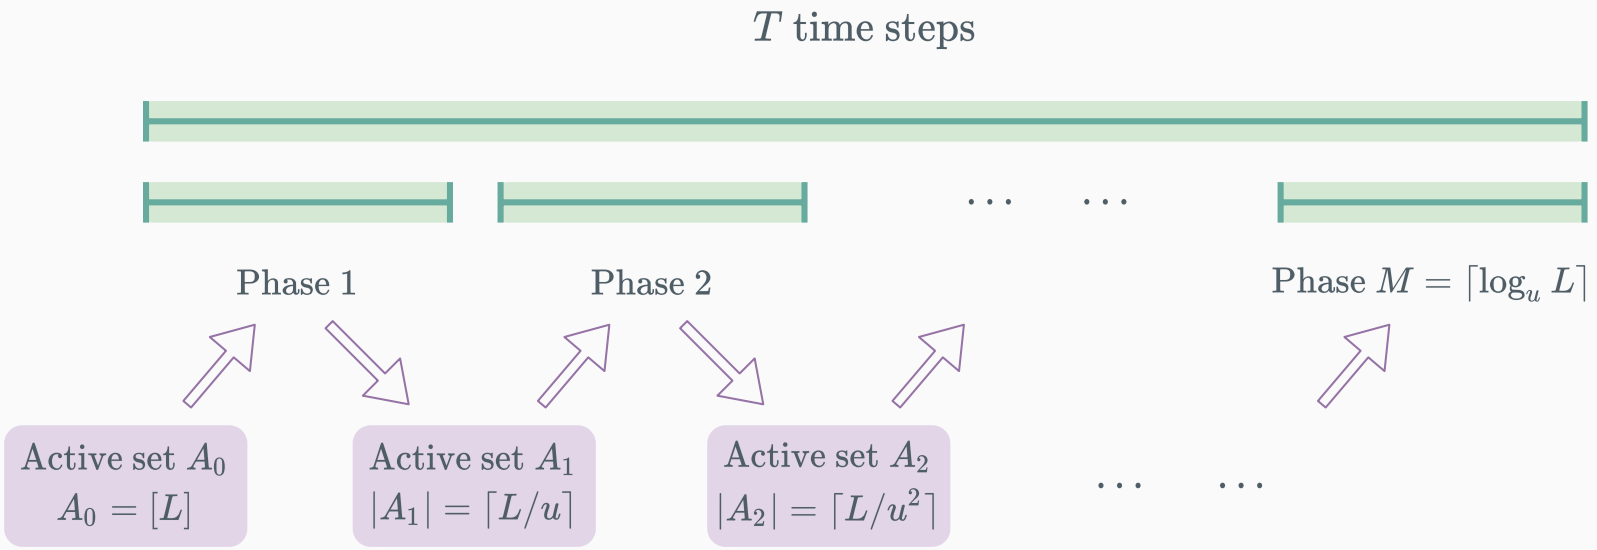

Probabilistic Sequential Shrinking **Algorithm**

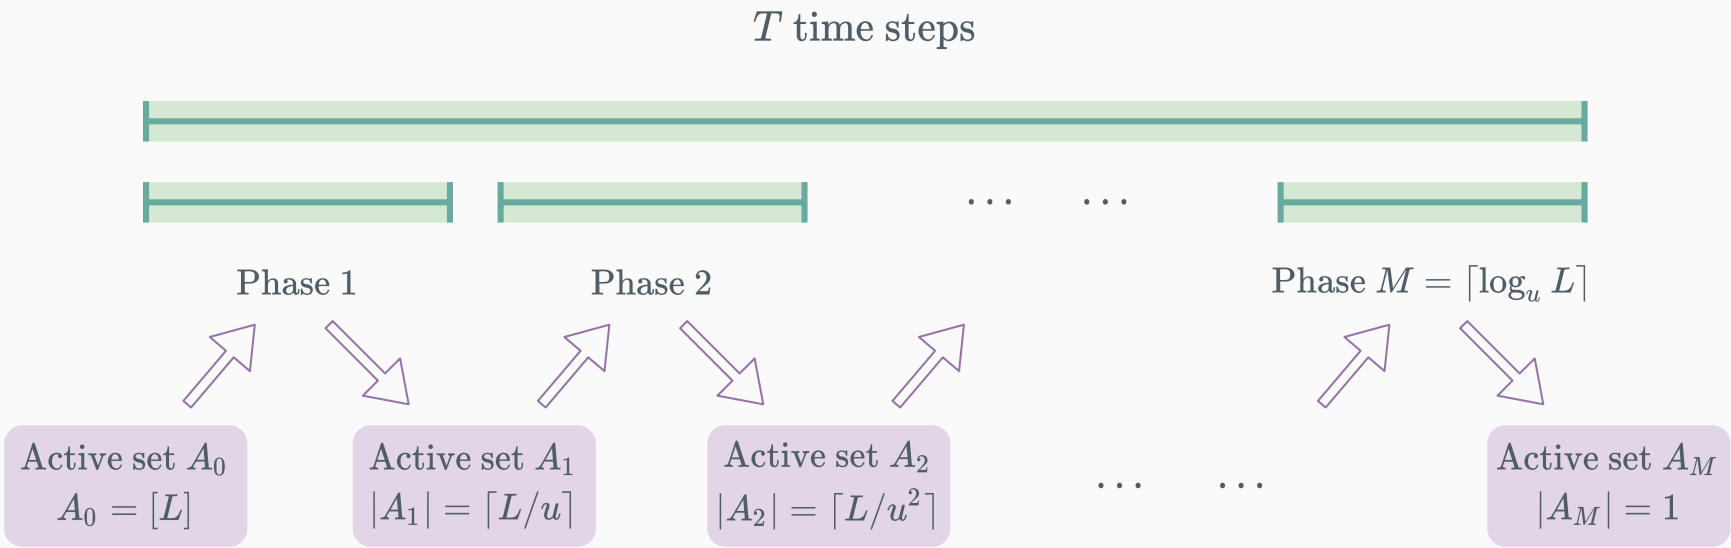

Probabilistic Sequential Shrinking **Algorithm**

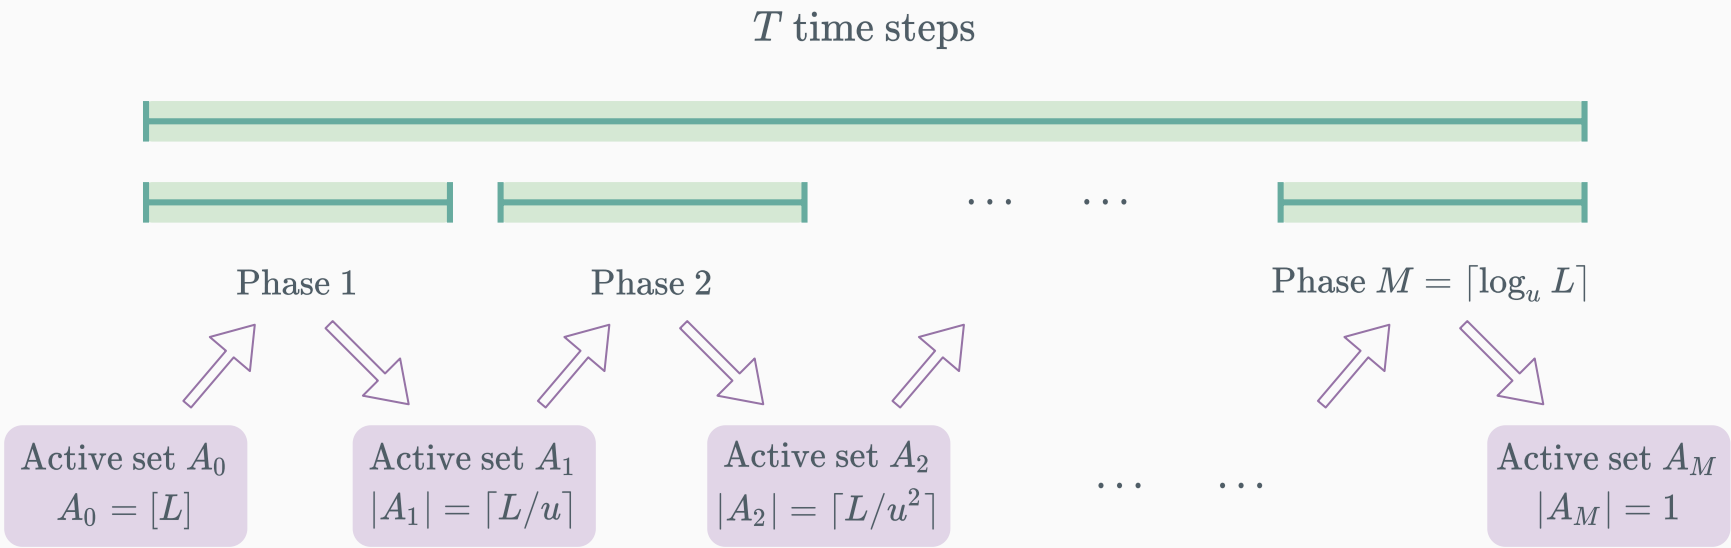

♠ How to shrink the active set?

PSS: **Shrink the active set**

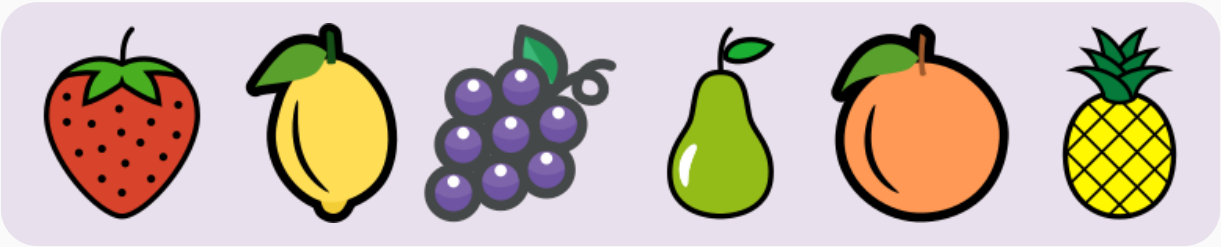

PSS: Shrink the active set

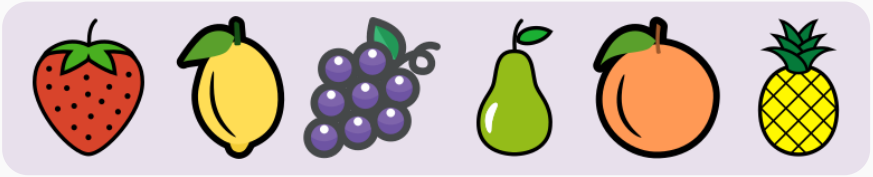

♠ Pull each active item with the same probability

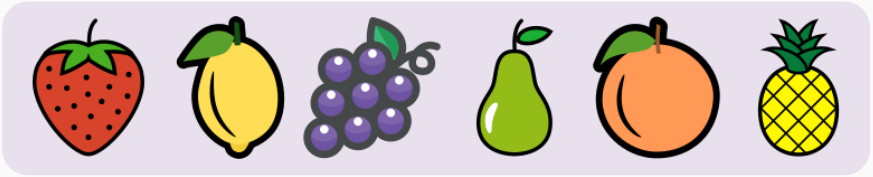

♠ Pull each active item with the same probability

PSS: Shrink the active set

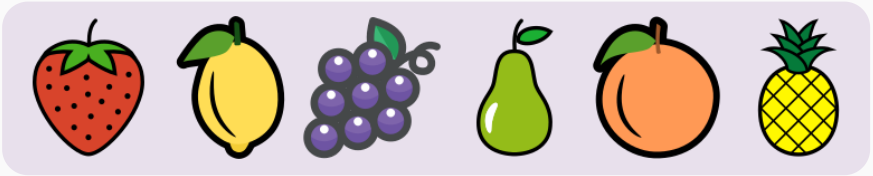

♠ Pull each active item with the same probability

$\tilde{W}_1(5) = 0.5$

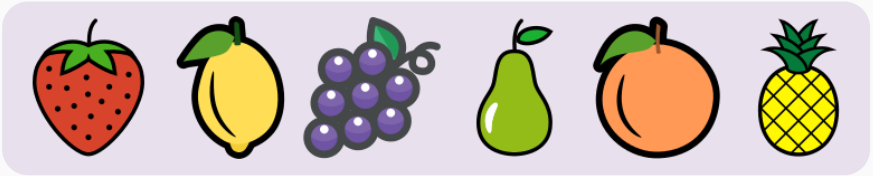

♠ Pull each active item with the same probability

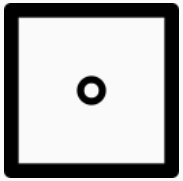

$\tilde{W}_1(5) = 0.5$

PSS: Shrink the active set

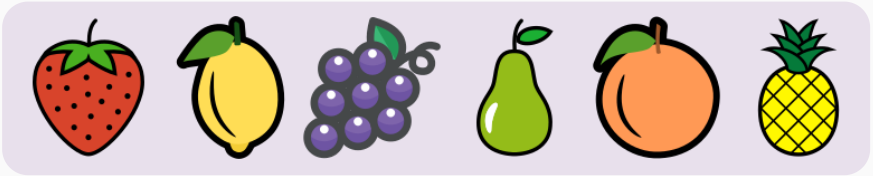

♠ Pull each active item with the same probability

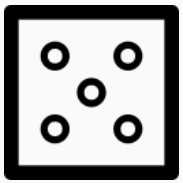

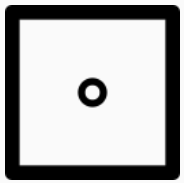

$\tilde{W}_1(5) = 0.5$ $\tilde{W}_2(1) = 0.3$

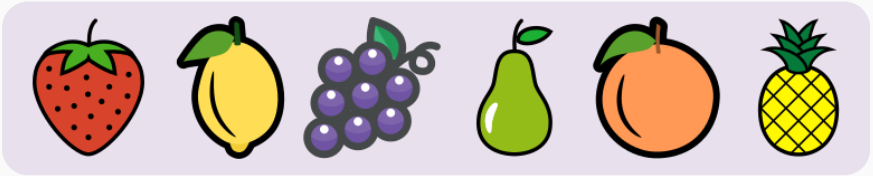

♠ Pull each active item with the same probability

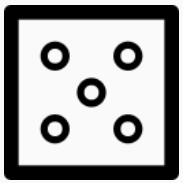

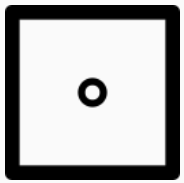

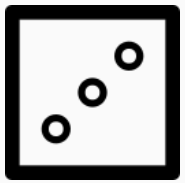

$\tilde{W}_1(5) = 0.5$ $\tilde{W}_2(1) = 0.3$

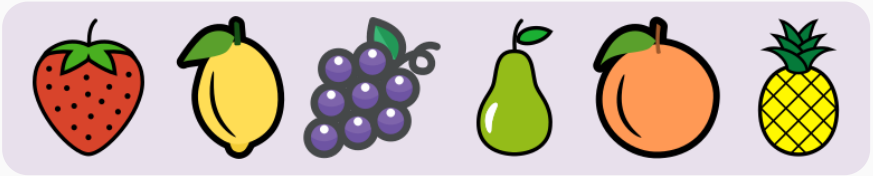

♠ Pull each active item with the same probability

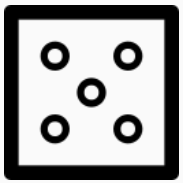

$\tilde{W}_1(5) = 0.5$

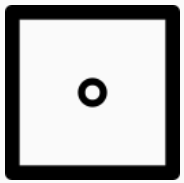

$\tilde{W}_2(1) = 0.3$

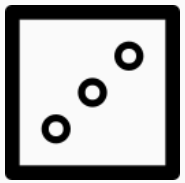

$\tilde{W}_3(3) = 0.9$

PSS: Shrink the active set

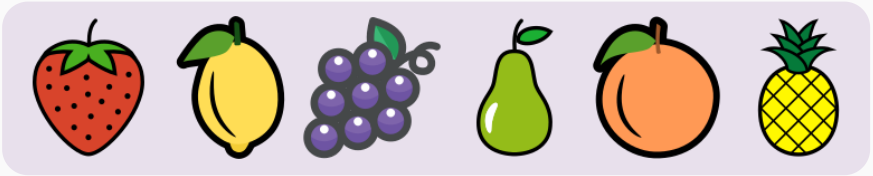

♠ Pull each active item with the same probability

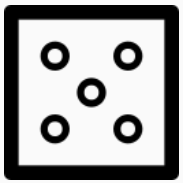

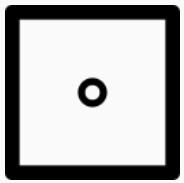

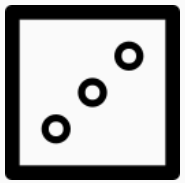

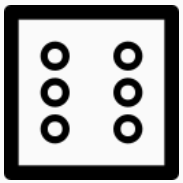

$\tilde{W}_1(5) = 0.5$ $\tilde{W}_2(1) = 0.3$ $\tilde{W}_3(3) = 0.9$

PSS: Shrink the active set

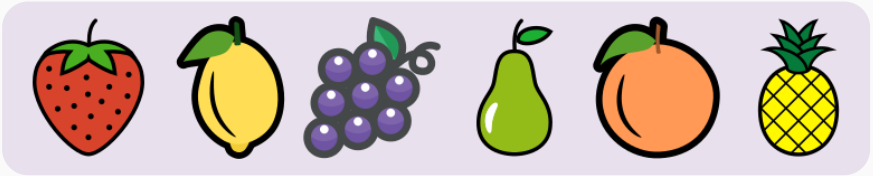

♠ Pull each active item with the same probability

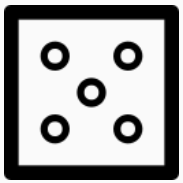

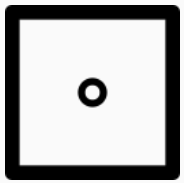

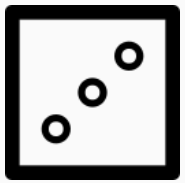

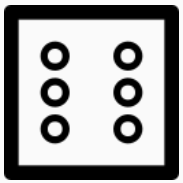

$\tilde{W}_1(5) = 0.5$ $\tilde{W}_2(1) = 0.3$ $\tilde{W}_3(3) = 0.9$ $\tilde{W}_4(6) = 0.2$

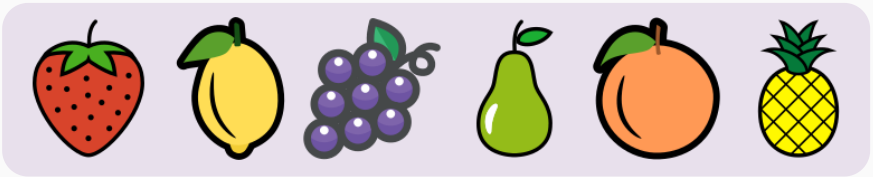

♠ Pull each active item with the same probability


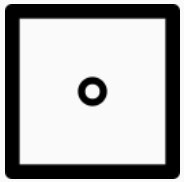

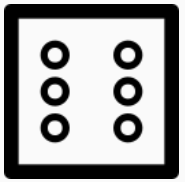

$\cdots$ $\cdots$

$\tilde{W}_1(5) = 0.5$ $\tilde{W}_2(1) = 0.3$ $\tilde{W}_3(3) = 0.9$ $\tilde{W}_4(6) = 0.2$ $\cdots$ $\cdots$

PSS: Shrink the active set

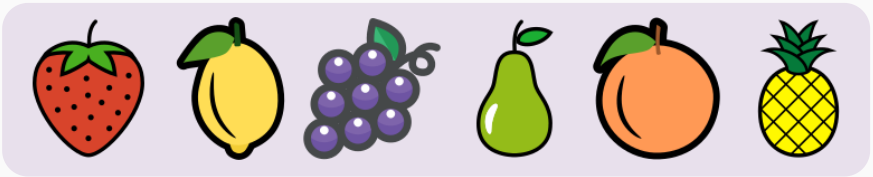

♠ Pull each active item with the same probability


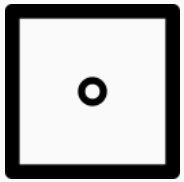

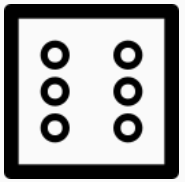

$\cdots$ $\cdots$

$\tilde{W}_1(5) = 0.5$ $\tilde{W}_2(1) = 0.3$ $\tilde{W}_3(3) = 0.9$ $\tilde{W}_4(6) = 0.2$ $\cdots$ $\cdots$

♠ Shrink the active set:
only keep items with high empirical means during the current phase

PSS: Shrink the active set

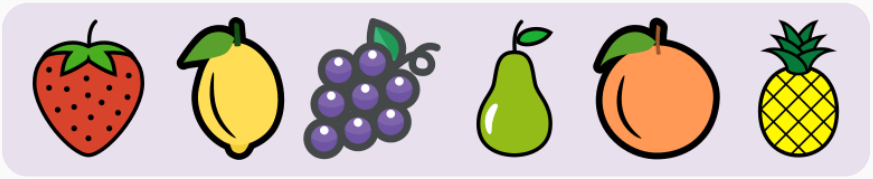

♠ Pull each active item with the same probability


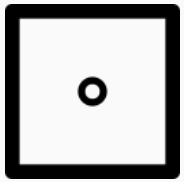

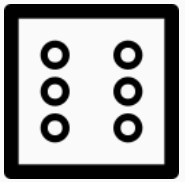

$\tilde{W}_1(5) = 0.5$ $\tilde{W}_2(1) = 0.3$ $\tilde{W}_3(3) = 0.9$ $\tilde{W}_4(6) = 0.2$ $\cdots$ $\cdots$

♠ Shrink the active set:
only keep items with high empirical means during the current phase

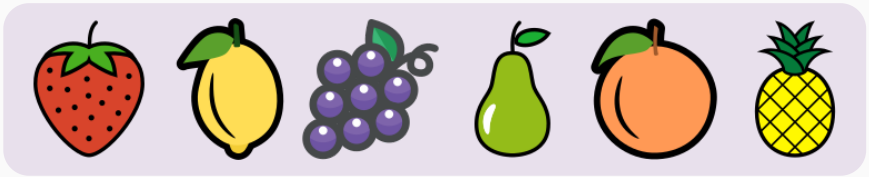

$\hat{W}_m(i)$	0.4	0.1	0.87	0.3	0.35	0.8

PSS: Shrink the active set

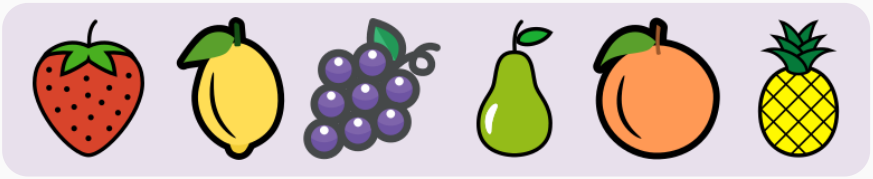

♠ Pull each active item with the same probability


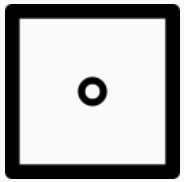

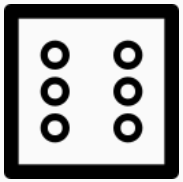

$\cdots \quad \cdots$

$\tilde{W}_1(5) = 0.5 \quad \tilde{W}_2(1) = 0.3 \quad \tilde{W}_3(3) = 0.9 \quad \tilde{W}_4(6) = 0.2 \quad \cdots \quad \cdots$

♠ Shrink the active set:
only keep items with high empirical means during the current phase

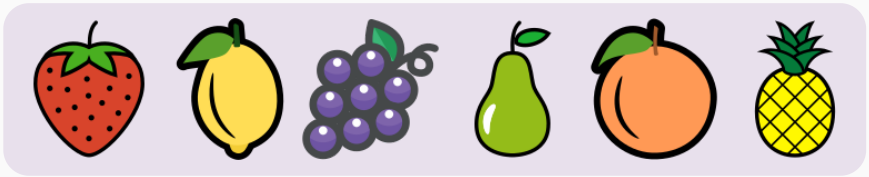

$\hat{W}_m(i)$ **0.4** 0.1 **0.87** 0.3 0.35 **0.8**

PSS: Shrink the active set

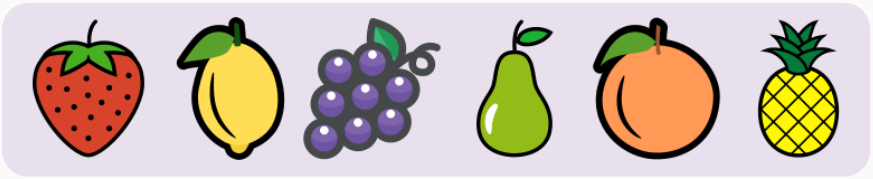

♠ Pull each active item with the same probability


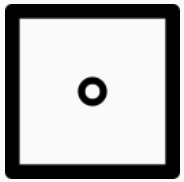

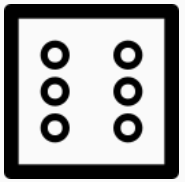

$\cdots \quad \cdots$

$\tilde{W}_1(5) = 0.5 \quad \tilde{W}_2(1) = 0.3 \quad \tilde{W}_3(3) = 0.9 \quad \tilde{W}_4(6) = 0.2 \quad \cdots \quad \cdots$

♠ Shrink the active set:
only keep items with high empirical means during the current phase

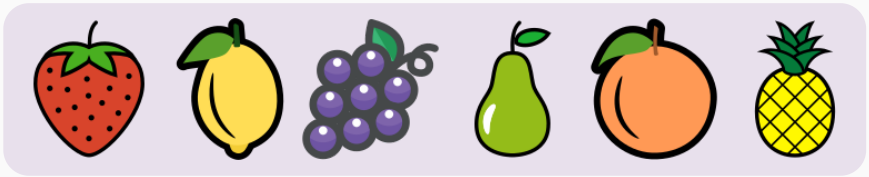

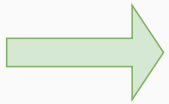

$\hat{W}_m(i)$ **0.4** 0.1 **0.87** 0.3 0.35 **0.8**

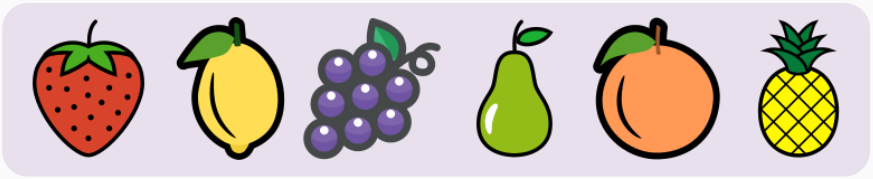

♠ Pull each active item with the same probability


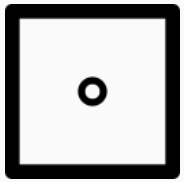

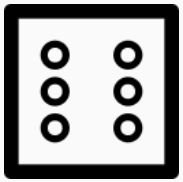

$\cdots \quad \cdots$

$\tilde{W}_1(5) = 0.5 \quad \tilde{W}_2(1) = 0.3 \quad \tilde{W}_3(3) = 0.9 \quad \tilde{W}_4(6) = 0.2 \quad \cdots \quad \cdots$

♠ Shrink the active set:
only keep items with high empirical means during the current phase

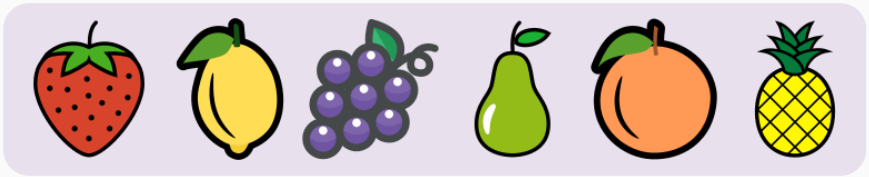

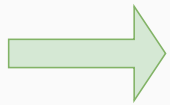

$\hat{W}_m(i)$ **0.4** 0.1 **0.87** 0.3 0.35 **0.8**

Comparison to deterministic algorithms: UP, SH

PSS(L) and Uniform Pull (UP)

- PSS(L): pulls each item for T/L times in expectation.
- UP: pulls each item for $\lfloor T/L \rfloor$ times with a deterministic schedule.

$\Rightarrow$ PSS(L): randomized version of UP.

Comparison to deterministic algorithms: UP, SH

PSS(L) and Uniform Pull (UP)

- PSS(L): pulls each item for T/L times in expectation.
- UP: pulls each item for $\lfloor T/L \rfloor$ times with a deterministic schedule.

$\Rightarrow$ PSS(L): randomized version of UP.

PSS(2) and Sequential Halving (SH) (Karnin et al., 2013)

- **Similarity**: both divide the whole horizon into $\lceil \log_2 L \rceil$ phases and halve the active set during each phase.
- **Difference**:
  - at each time step of phase m, PSS(2) chooses item $i \in A_{m-1}$ with probability $1/|A_{m-1}|$ and pulls it;
  - during phase m, SH pulls each item in A_{m-1} for exactly $\lfloor T/(\lceil \log_2 L \rceil \cdot |A_{m-1}|) \rfloor$ times according to a deterministic schedule.

$\Rightarrow$ PSS(2): randomized version of SH.

Comparison among upper bounds

Comparison in stochastic bandits with adversarial corruptions

Algorithm	Order of error bound ϵ_C	Order of failure probability δ
PSS(u)	$\frac{C \log_u L}{T}$	$L(\log_u L) \exp\left[- \frac{T}{192 \tilde{H}_2(w, L, u) \log_u L} \right]$

Comparison among upper bounds

Comparison in stochastic bandits with adversarial corruptions

Algorithm	Order of error bound ϵ_C	Order of failure probability δ
PSS(u)	$\frac{C \log_u L}{T}$	$L(\log_u L) \exp\left[- \frac{T}{192 \tilde{H}_2(w, L, u) \log_u L} \right]$
PSS(2)	$\frac{C \log_2 L}{T}$	$L(\log_2 L) \exp\left[- \frac{T}{192 \tilde{H}_2(w, L, u) \log_2 L} \right]$
SH	$\frac{CL \log_2 L}{T}$	$L(\log_2 L) \exp\left[- \frac{T}{192 \tilde{H}_2(w, L, u) \log_2 L} \right]$

Comparison among upper bounds

Comparison in stochastic bandits with adversarial corruptions

Algorithm	Order of error bound ϵ_C	Order of failure probability δ
PSS(u)	$\frac{C \log_u L}{T}$	$L(\log_u L) \exp\left[- \frac{T}{192 \tilde{H}_2(w, L, u) \log_u L} \right]$
PSS(2)	$\frac{C \log_2 L}{T}$	$L(\log_2 L) \exp\left[- \frac{T}{192 \tilde{H}_2(w, L, u) \log_2 L} \right]$
SH	$\frac{CL \log_2 L}{T}$	$L(\log_2 L) \exp\left[- \frac{T}{192 \tilde{H}_2(w, L, u) \log_2 L} \right]$
PSS(L)	$\frac{C}{T}$	$L \exp\left(- \frac{T}{192 L / \Delta_{1,2}^2} \right)$
UP	$\frac{CL}{T}$	$L \exp\left(- \frac{T}{192 L / \Delta_{1,2}^2} \right)$

- $\tilde{H}_2(w, L, u) = \max\limits_{i \neq 1} \frac{\min\{u \cdot i,\ L\}}{\Delta_{1,i}^2}$: quantify difficulty of BAI.
- $H_2(w) = \max\limits_{i \neq 1} \frac{i}{\Delta_i^2}$, $\tilde{H}_2(w, L, 1) = H_2(w)$, $\tilde{H}_2(w, L, u) \leq u \cdot H_2(w)$.

Theorem 2.2

Fix $\lambda \in (0,1)$ and $\Delta \in (0,1/2)$. For any online algorithm, there is a BAI with an adversarial corruption instance over T steps, corruption budget $C = 1 + (1+\lambda)2\Delta T$, and optimality gap Δ, such that

$$\begin{aligned}\mathbb{P}[\Delta_{1,i_{\text{out}}} > 0] = \mathbb{P}[\Delta_{1,i_{\text{out}}} \geq \Delta] &= \mathbb{P}[i_{\text{out}} \neq 1] \\ &\geq \frac{1}{2} \cdot \left[1 - \exp\left(-\frac{2\lambda^2 \Delta T}{3}\right)\right].\end{aligned}$$

- $\frac{C}{T} > 2\Delta_{1,2}$: It is impossible for any algorithm to identify the optimal item with high probability.
- $\frac{C}{T} \leq \frac{\Delta_{1,L}}{8\lceil \log_u L \rceil}$: our work (Theorem 4.1) provides a guarantee for PSS(u).

$\Rightarrow$ The upper bound in our work (Theorem 4.1) is within a factor of $O(\log L)$ away from the largest possible upper bound on C/T in Theorem **??**.

Outline

1. What is multi-armed bandits (MAB)?

Summary

- Introduction on multi-armed bandit problems
- Problem formulation
    - Hardness H_1, H_2; concentration inequalities

Summary

- Introduction on multi-armed bandit problems
- Problem formulation
  - Hardness H_1, H_2; concentration inequalities
- BAI under the fixed-confidence setting
  - Algorithms: Successive elimination, Median Elimination, Track & Stop
  - Lower bound: achieved by Track & Stop

Summary

- Introduction on multi-armed bandit problems
- Problem formulation
  - Hardness H_1, H_2; concentration inequalities
- BAI under the fixed-confidence setting
  - Algorithms: Successive elimination, Median Elimination, Track & Stop
  - Lower bound: achieved by Track & Stop
- BAI under the fixed-budget setting
  - Algorithms: UCB-E, Sequential Halving
  - Gap between upper and lower bounds
  - With adversarial corruptions: Probabilistic Sequential Shrinking

Summary

- Introduction on multi-armed bandit problems
- Problem formulation
  - Hardness H_1, H_2; concentration inequalities
- BAI under the fixed-confidence setting
  - Algorithms: Successive elimination, Median Elimination, Track & Stop
  - Lower bound: achieved by Track & Stop
- BAI under the fixed-budget setting
  - Algorithms: UCB-E, Sequential Halving
  - Gap between upper and lower bounds
  - With adversarial corruptions: Probabilistic Sequential Shrinking
- More existing works ...
  - *Multiple pure exploration*: to identify multiple arms
    CLUCB by Chen et al. (2014), EST1 and CSAR by Rejwan and Mansour (2020)
  - *Pure exploration in linear bandits*
    (Jedra and Proutiere; Yang and Tan, 2021)
  - • • •

Further exploration

- Fill the gap between upper and lower bounds for BAI under the fixed-budget setting?

Further exploration

- Fill the gap between upper and lower bounds for BAI under the fixed-budget setting?
- Identification of the arm with the highest median reward (Altschuler et al., 2019):
  More studies taking the median of rewards as the criterion are yet to be done.

Further exploration

- Fill the gap between upper and lower bounds for BAI under the fixed-budget setting?
- Identification of the arm with the highest median reward (Altschuler et al., 2019):
  More studies taking the median of rewards as the criterion are yet to be done.
- BAI in adversarial bandits (Shen, 2019; Zhong et al., 2021):
  Optimal attack strategies against regret minimization (Jun et al., 2018; Liu and Lai, 2020)
  Optimal attack strategies against pure exploration?

Thanks for listening!

https://zixinzh.github.io/homepage/conf_tutorial/

Emails: vtan@nus.edu.sg, zixin.zhong@u.nus.edu

References I

N. Abe and P. M. Long. Associative reinforcement learning using linear probabilistic concepts. In *Proceedings of the 16th International Conference on Machine Learning*, pages 3–11, 1999.

M. Abramowitz and I. A. Stegun. *Handbook of Mathematical Functions with Formulas, Graphs, and Mathematical Tables*. Dover, New York, ninth edition, 1964.

S. Agrawal and N. Goyal. Analysis of Thompson sampling for the multi-armed bandit problem. In *Proceedings of the 25th Annual Conference on Learning Theory*, pages 39.1–39.26, 2012.

J. Altschuler, V.-E. Brunel, and A. Malek. Best arm identification for contaminated bandits. *Journal of Machine Learning Research*, 20(91):1–39, 2019.

J.-Y. Audibert and S. Bubeck. Best arm identification in multi-armed bandits. In *Proceedings of the 23th Conference on Learning Theory*, pages 41–53, 2010.

J.-Y. Audibert, S. Bubeck, et al. Minimax policies for adversarial and stochastic bandits. In *Proceedings of the 22th Conference on Learning Theory*, pages 1–122, 2009.

References II

P. Auer, N. Cesa-Bianchi, and P. Fischer. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47(2-3):235–256, 2002a.

P. Auer, N. Cesa-Bianchi, Y. Freund, and R. E. Schapire. The nonstochastic multiarmed bandit problem. *SIAM Journal of Computing*, 32(1):48–77, 2002b.

D. A. Berry, R. W. Chen, A. Zame, D. C. Heath, and L. A. Shepp. Bandit problems with infinitely many arms. *The Annals of Statistics*, pages 2103–2116, 1997.

S. Bubeck, T. Wang, and N. Viswanathan. Multiple identifications in multi-armed bandits. In *Proceedings of the 30th International Conference on Machine Learning*, pages 258–265, 2013.

A. Carpentier and A. Locatelli. Tight (lower) bounds for the fixed budget best arm identification bandit problem. In *Proceedings of the 29th Conference on Learning Theory*, pages 590–604, 2016.

A. Carpentier and M. Valko. Simple regret for infinitely many armed bandits. In *Proceedings of the 32nd International Conference on Machine Learning*, pages 1133–1141, 2015.

N. Cesa-Bianchi and G. Lugosi. *Prediction, learning, and games*. Cambridge university press, 2006.

References III

S. Chen, T. Lin, I. King, M. R. Lyu, and W. Chen. Combinatorial pure exploration of multi-armed bandits. In *Proceedings of the 27th Advances in Neural Information Processing Systems*, pages 379–387. 2014.

D. P. Dubhashi and A. Panconesi. *Concentration of measure for the analysis of randomized algorithms*. Cambridge University Press, 2009.

E. Even-Dar, S. Mannor, and Y. Mansour. PAC bounds for multi-armed bandit and markov decision processes. In *Proceedings of the 15th Annual International Conference on Computational Learning Theory*, pages 255–270, 2002.

V. Gabillon, M. Ghavamzadeh, and A. Lazaric. Best arm identification: A unified approach to fixed budget and fixed confidence. In *Proceedings of the 25th International Conference on Neural Information Processing Systems*, 2012.

A. Garivier and E. Kaufmann. Optimal best arm identification with fixed confidence. In V. Feldman, A. Rakhlin, and O. Shamir, editors, *29th Annual Conference on Learning Theory*, volume 49 of *Proceedings of Machine Learning Research*, pages 998–1027, Columbia University, New York, New York, USA, 23–26 Jun 2016. PMLR.

References IV

A. Gupta, T. Koren, and K. Talwar. Better algorithms for stochastic bandits with adversarial corruptions. In *Proceedings of the 32nd Conference on Learning Theory*, pages 1562–1578, 2019.

K. Jamieson and R. Nowak. Best-arm identification algorithms for multi-armed bandits in the fixed confidence setting. In *Proceedings of the 48th Annual Conference on Information Sciences and Systems (CISS)*, pages 1–6, 2014.

Y. Jedra and A. Proutiere. Optimal best-arm identification in linear bandits. In H. Larochelle, M. Ranzato, R. Hadsell, M. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, pages 10007–10017. Curran Associates, Inc.

K.-S. Jun, L. Li, Y. Ma, and J. Zhu. Adversarial attacks on stochastic bandits. In *Proceedings of the 31st Advances in Neural Information Processing Systems*, pages 3640–3649, 2018.

S. Kalyanakrishnan, A. Tewari, P. Auer, and P. Stone. Pac subset selection in stochastic multi-armed bandits. In *Proceedings of the 29th International Conference on Machine Learning*, pages 655–662, 2012.

References V

Z. Karnin, T. Koren, and O. Somekh. Almost optimal exploration in multi-armed bandits. In *Proceedings of the 13th International Conference on Machine Learning*, pages 1238–1246, 2013.

Y. Kuroki, L. Xu, A. Miyauchi, J. Honda, and M. Sugiyama. Polynomial-time algorithms for multiple-arm identification with full-bandit feedback. *Neural Computation*, 32(9):1733–1773, 2020.

B. Kveton, C. Szepesvari, Z. Wen, and A. Ashkan. Cascading bandits: Learning to rank in the cascade model. In *Proceedings of the 32nd International Conference on Machine Learning*, pages 767–776, 2015a.

B. Kveton, Z. Wen, A. Ashkan, and C. Szepesvári. Combinatorial cascading bandits. In *Proceedings of the 28th International Conference on Neural Information Processing Systems*, pages 1450–1458, 2015b.

T. L. Lai. Adaptive treatment allocation and the multi-armed bandit problem. *The Annals of Statistics*, 15(3):1091 – 1114, 1987.

S. Li, B. Wang, S. Zhang, and W. Chen. Contextual combinatorial cascading bandits. In *Proceedings of the 33rd International Conference on Machine Learning*, pages 1245–1253, 2016.

References VI

G. Liu and L. Lai. Action-manipulation attacks on stochastic bandits. In *Proceedings of the 45th International Conference on Acoustics, Speech and Signal Processing*, pages 3112–3116, 2020.

S. Mannor and J. N. Tsitsiklis. The sample complexity of exploration in the multi-armed bandit problem. *Journal of Machine Learning Research*, 5(Jun): 623–648, 2004.

I. Rejwan and Y. Mansour. Top-k combinatorial bandits with full-bandit feedback. In *Proceedings of the 31st International Conference on Algorithmic Learning Theory*, pages 752–776, 2020.

D. Russo and B. Van Roy. Learning to optimize via posterior sampling. *Mathematics of Operations Research*, 39(4):1221–1243, 2014.

S. Shahrampour, M. Noshad, and V. Tarokh. On sequential elimination algorithms for best-arm identification in multi-armed bandits. *IEEE Transactions on Signal Processing*, 65(16):4281–4292, 2017. doi: $10.1109/\mathrm{TSP}.2017.2706192$.

C. Shen. Universal best arm identification. *IEEE Transactions on Signal Processing*, 67(17):4464–4478, 2019.

J. Yang and V. Tan. Towards minimax optimal best arm identification in linear bandits. *arXiv preprint arXiv:2105.13017*, 2021.

References VII

Z. Zhong, W. C. Cheung, and V. Tan. Probabilistic sequential shrinking: A best arm identification algorithm for stochastic bandits with corruptions. In *Proceedings of the 38th International Conference on Machine Learning*, 2021.

J. Zimmert and Y. Seldin. Tsallis-INF: An optimal algorithm for stochastic and adversarial bandits. *Journal of Machine Learning Research*, 22:28–1, 2021.